Machine Learning Theory

Least Squares with Sampling, Misspecification, and Non-Gaussian Noise

David A. Hirshberg

April 2, 2025

Emory University

Where We Left Off

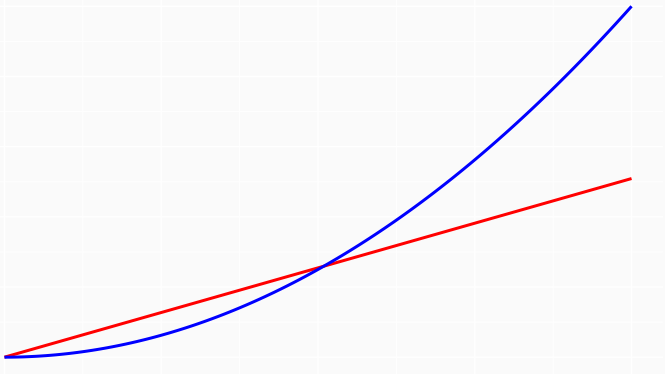

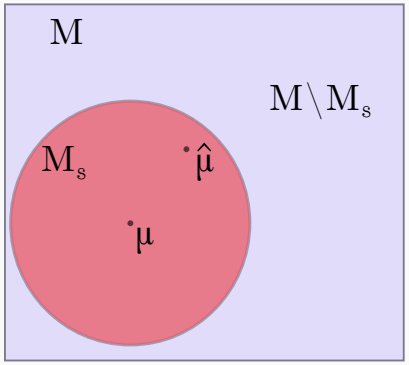

What do we know about the error of this least squares estimator $\hat{\mu}$?

$$\hat{\mu} = \underset{m \in \mathcal{M}}{\operatorname{argmin}} \frac{1}{n} \sum_{i=1}^{n} \{ Y_i - m(X_i) \}^2 \quad \text{for} \quad \text{convex} \ \mathcal{M}$$

Where We Left Off

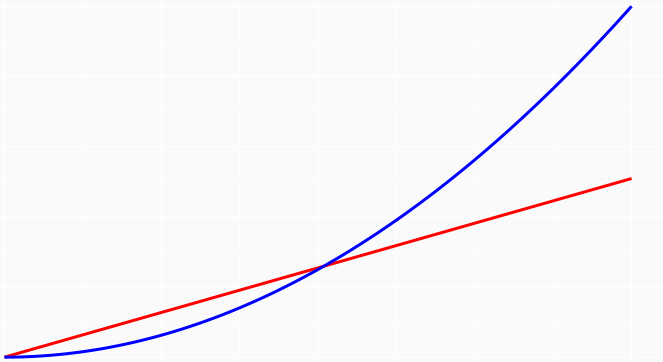

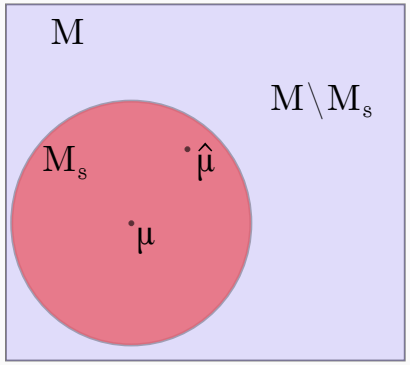

What do we know about the error of this least squares estimator $\hat{\mu}$?

$$\hat{\mu} = \underset{m \in \mathcal{M}}{\operatorname{argmin}} \frac{1}{n} \sum_{i=1}^{n} \{ Y_i - m(X_i) \}^2 \quad \text{for} \quad \text{convex} \ \mathcal{M}$$

Here's what we proved in lecture.

$$\|\hat{\mu} - \mu^\star\|_{L_2(\mathrm{P_n})} < s \quad \text{w.p.} \ 1 - \delta \ \text{for} \quad \frac{s^2}{2\sigma} \geq \mathrm{w}(\mathcal{M}_s^\circ) + s\sqrt{\frac{2\Sigma_n}{\delta n}}$$

where $\mathrm{w}(\mathcal{V}) = \mathrm{E} \max_{v \in \mathcal{V}} \langle g, \ v \rangle_{L_2(\mathrm{P_n})}$ and $\Sigma_n = \sigma^2 \{1 + 2\log(2n)\}$ for $g_i \overset{iid}{\sim} N(0,1)$

if $Y_i = \mu(X_i) + \varepsilon_i$ for $\varepsilon_i \overset{iid}{\sim} N(0, \sigma^2)$ for $\mu \in \mathcal{M}$

Where We Left Off

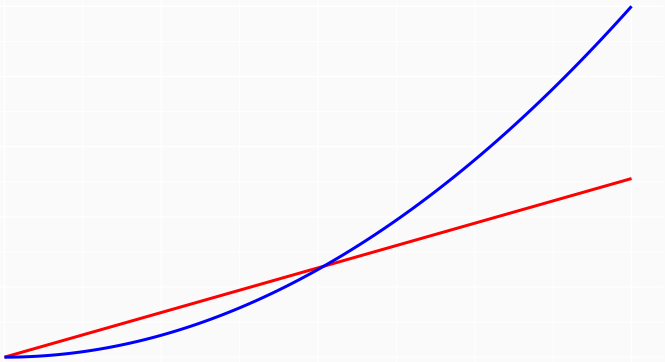

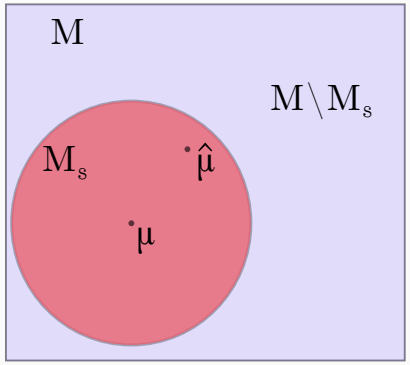

What do we know about the error of this least squares estimator $\hat{\mu}$?

$$\hat{\mu} = \underset{m \in \mathcal{M}}{\operatorname{argmin}} \frac{1}{n} \sum_{i=1}^{n} \{ Y_i - m(X_i) \}^2 \quad \text{for} \quad \text{convex} \ \mathcal{M}$$

Here's a simplified version of you're proving for homework.

$$\|\hat{\mu} - \mu^\star\|_{L_2(\mathrm{P_n})} < s + 2\sqrt{\frac{2\Sigma_n}{\delta n}} \quad \text{w.p.} \ \ 1 - \delta \ \ \text{for} \quad \frac{s^2}{2\sigma} \geq \mathrm{w}(\mathcal{M}_s)$$

where $\mathrm{w}(\mathcal{V}) = \mathrm{E} \max_{v \in \mathcal{V}} \langle g, \ v \rangle_{L_2(\mathrm{P_n})}$ and $\Sigma_n = \sigma^2 \{1 + 2\log(2n)\}$ for $g_i \overset{iid}{\sim} N(0,1)$

if $Y_i = \mu(X_i) + \varepsilon_i$ for $\varepsilon_i \overset{iid}{\sim} N(0, \sigma^2)$ for $\mu \in \mathcal{M}$

When Does This Bound Apply?

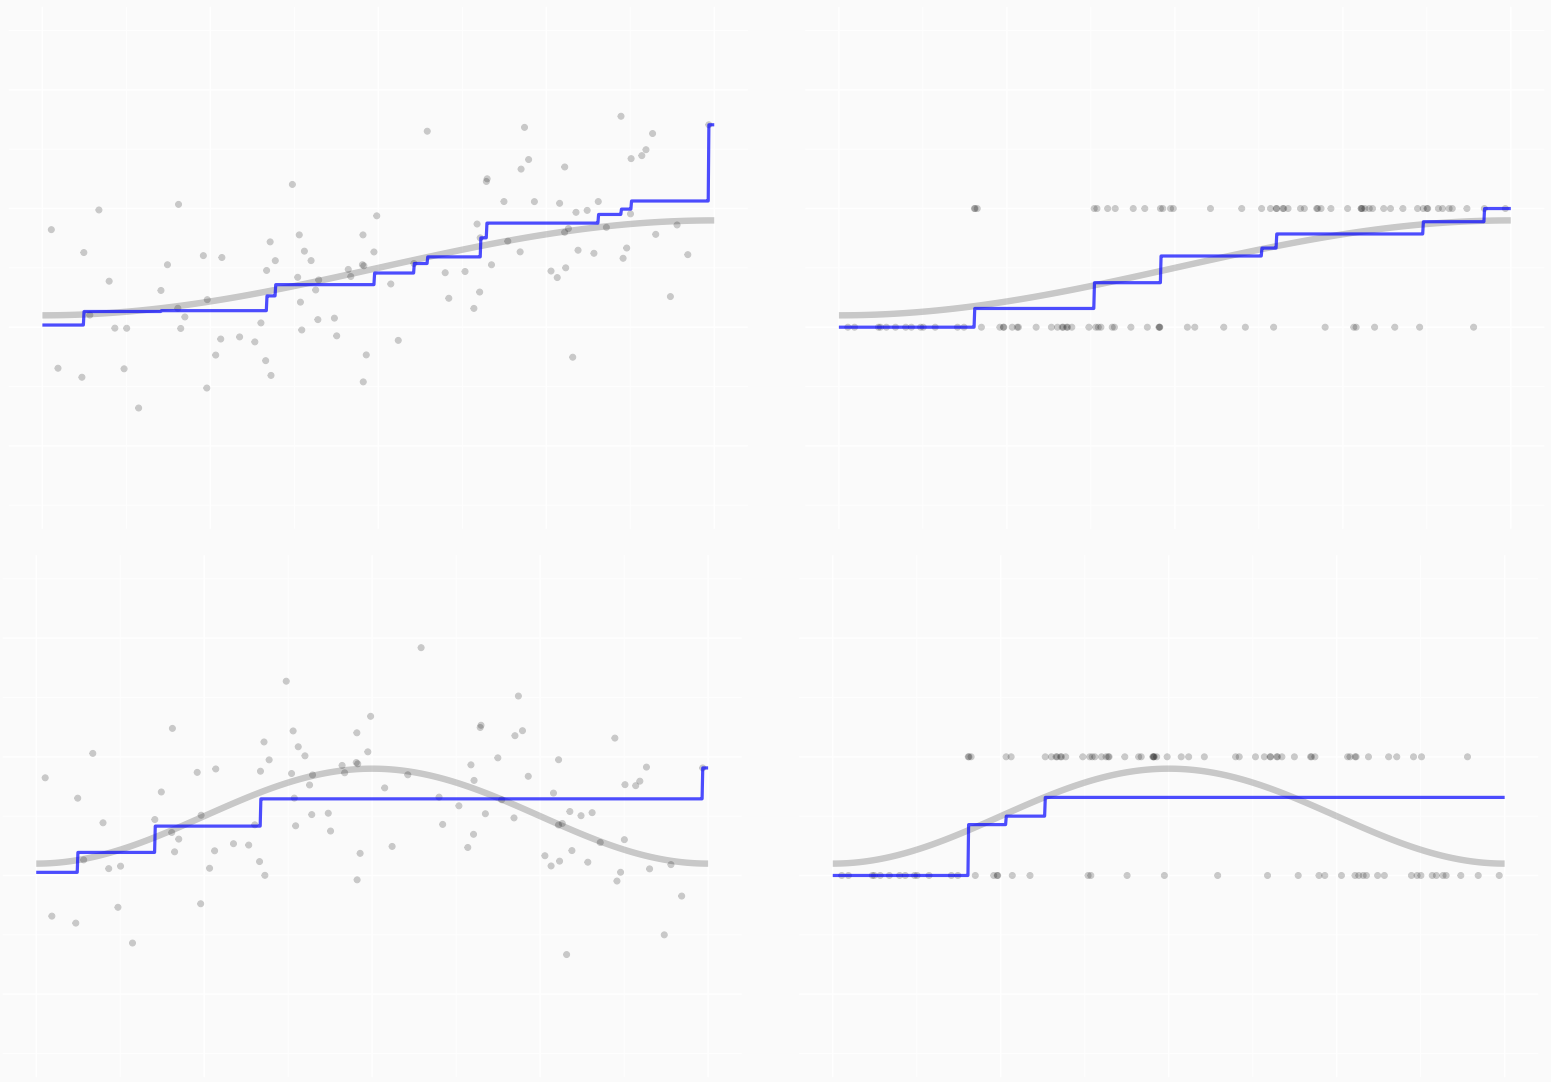

When Does This Bound Apply?

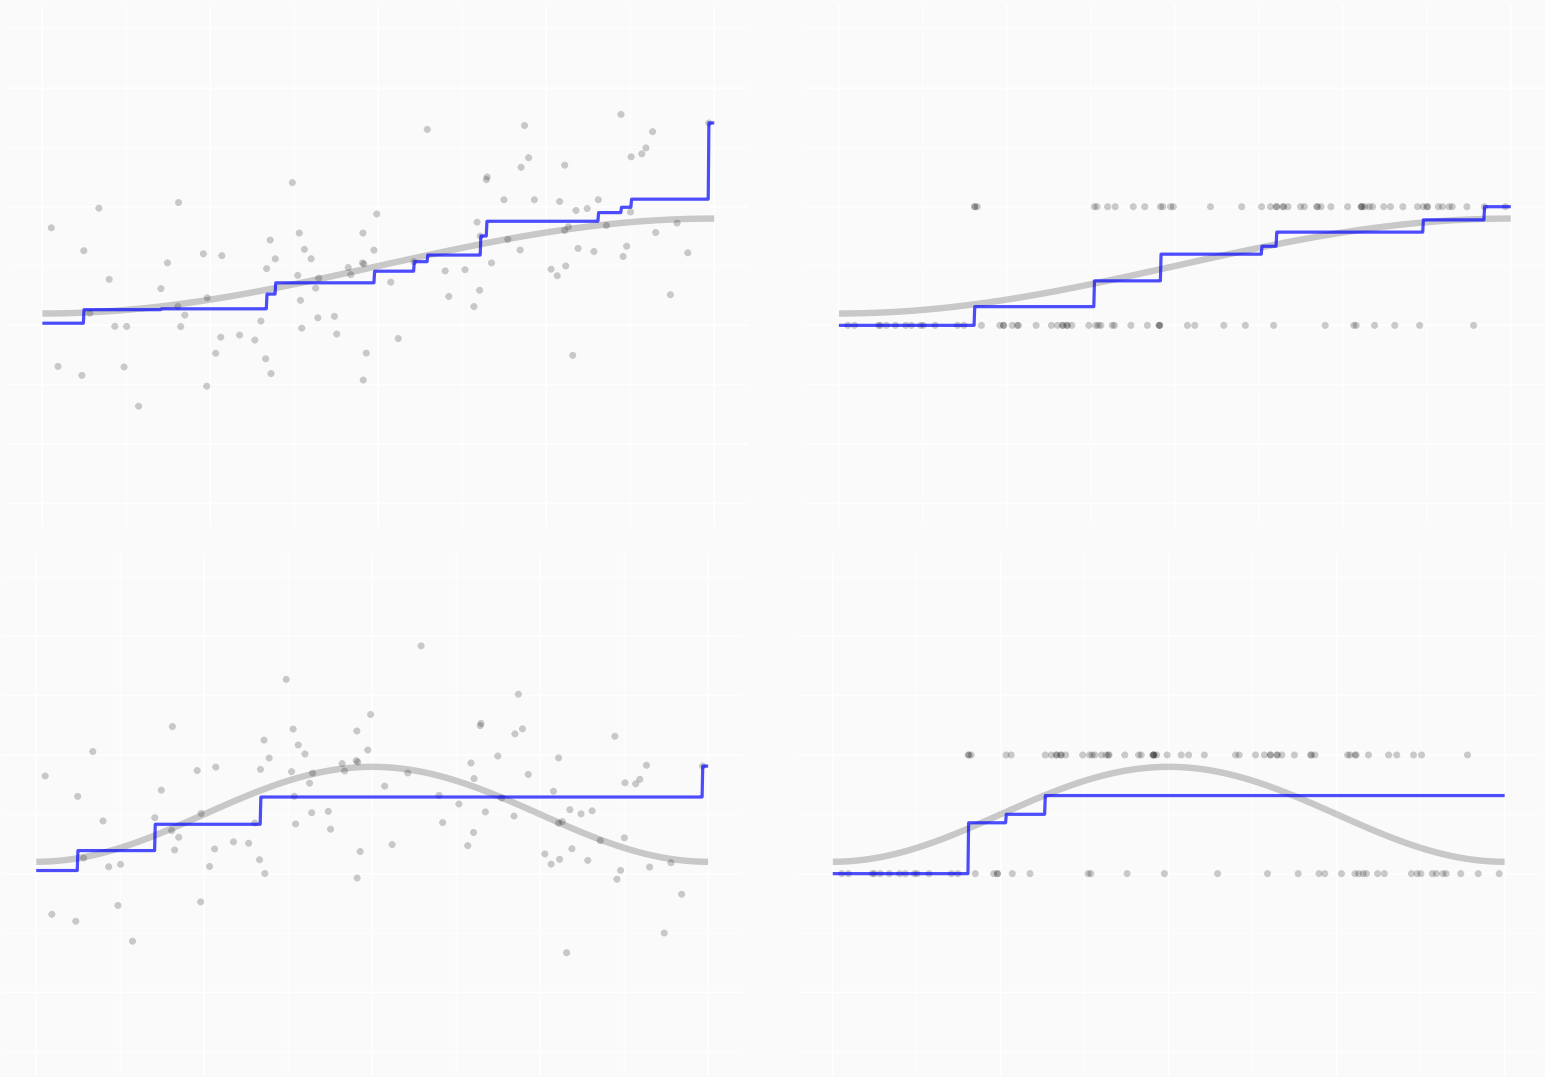

- The second column is out. We've assumed correct specfication.
- The second row is out. We've assumed normality.

Today, We Fix That

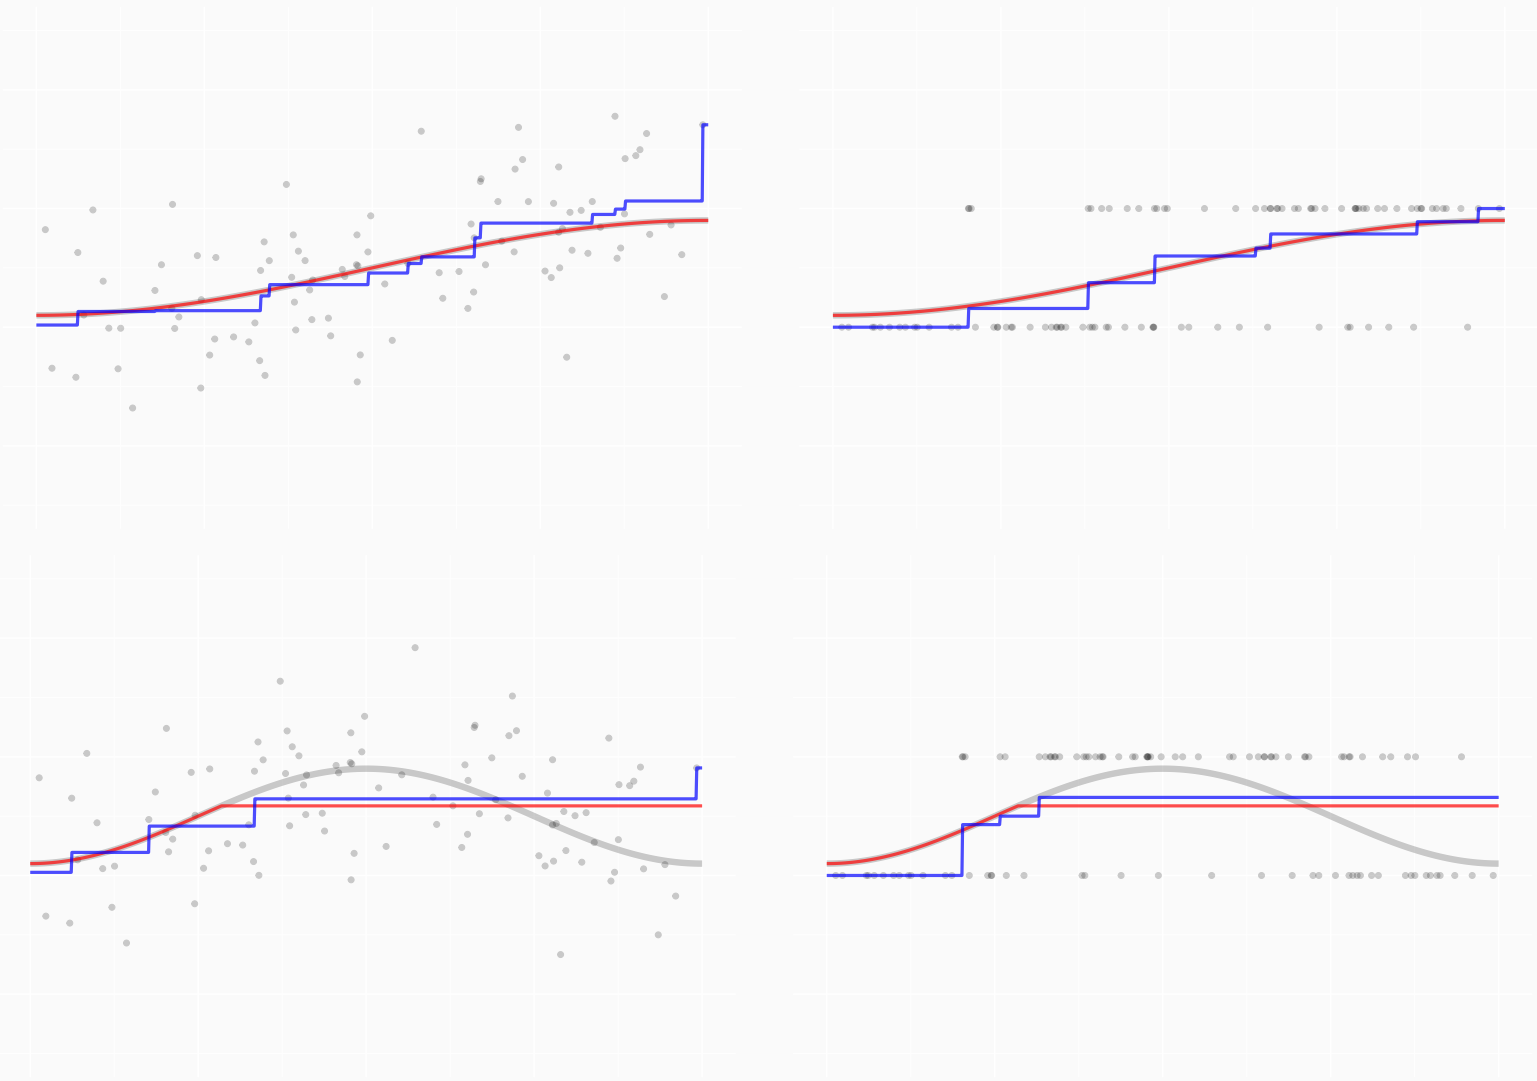

Today, We Fix That

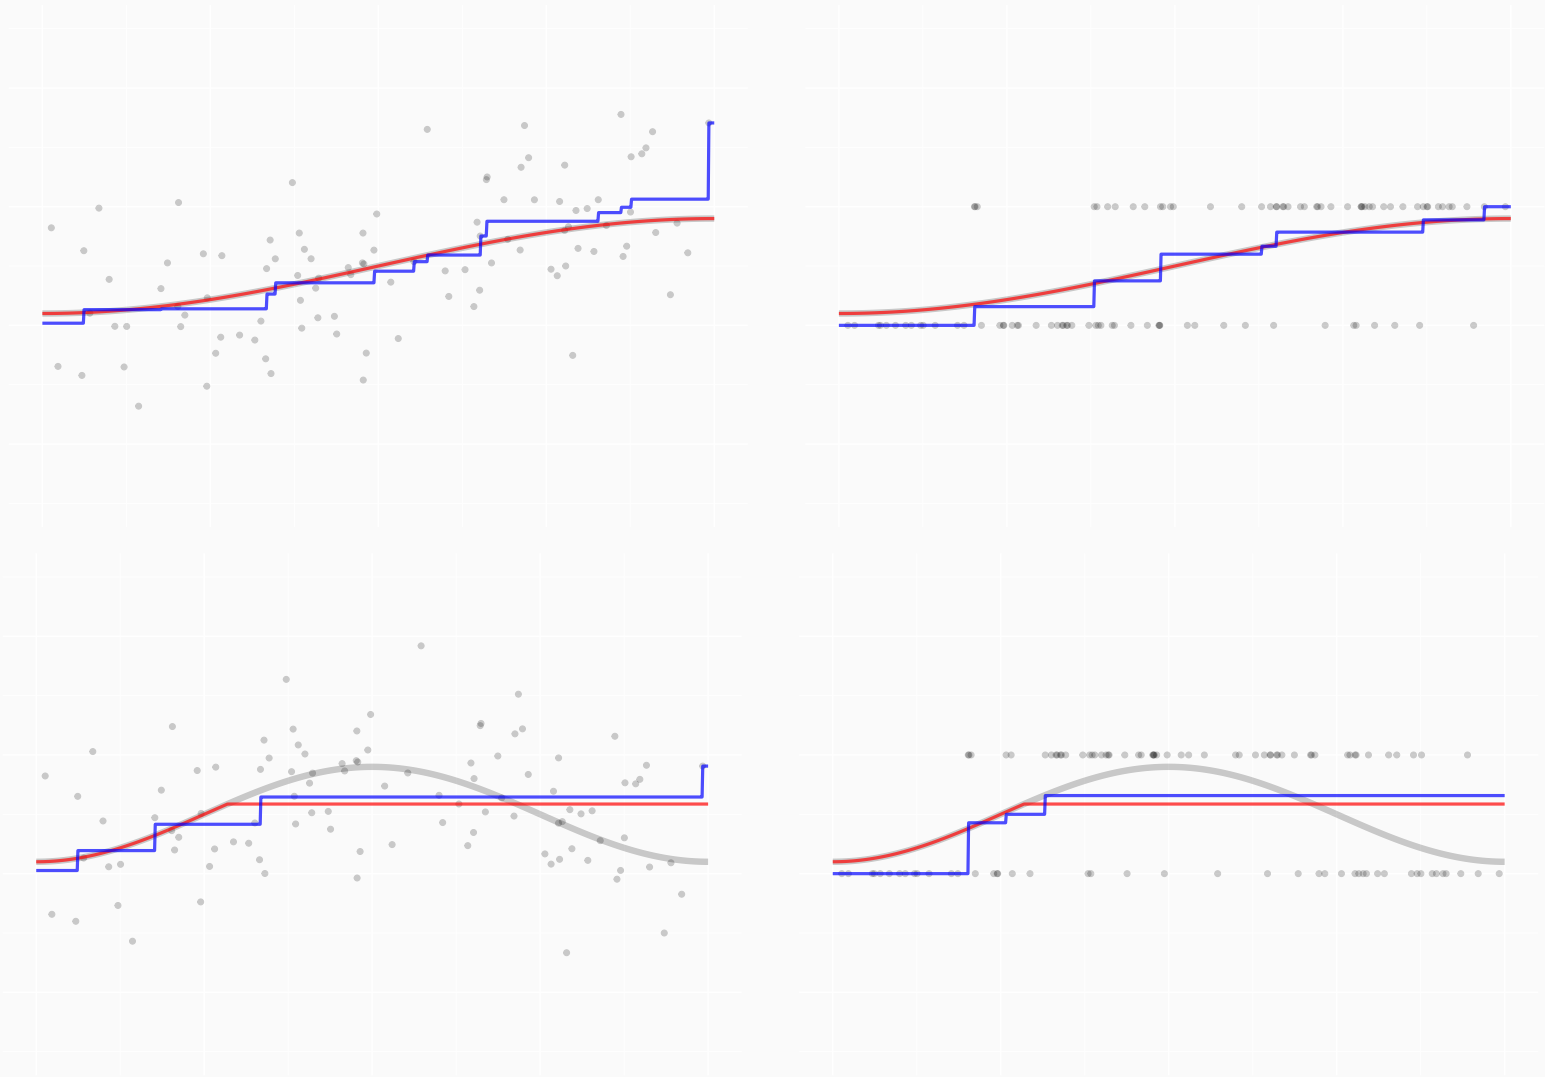

- With misspecification, we estimate the model's best approximation to μ.
- Non-normality doesn't really matter much. We'll look at how it affects our bound.

Misspecification

What happens when μ isn't in the model?

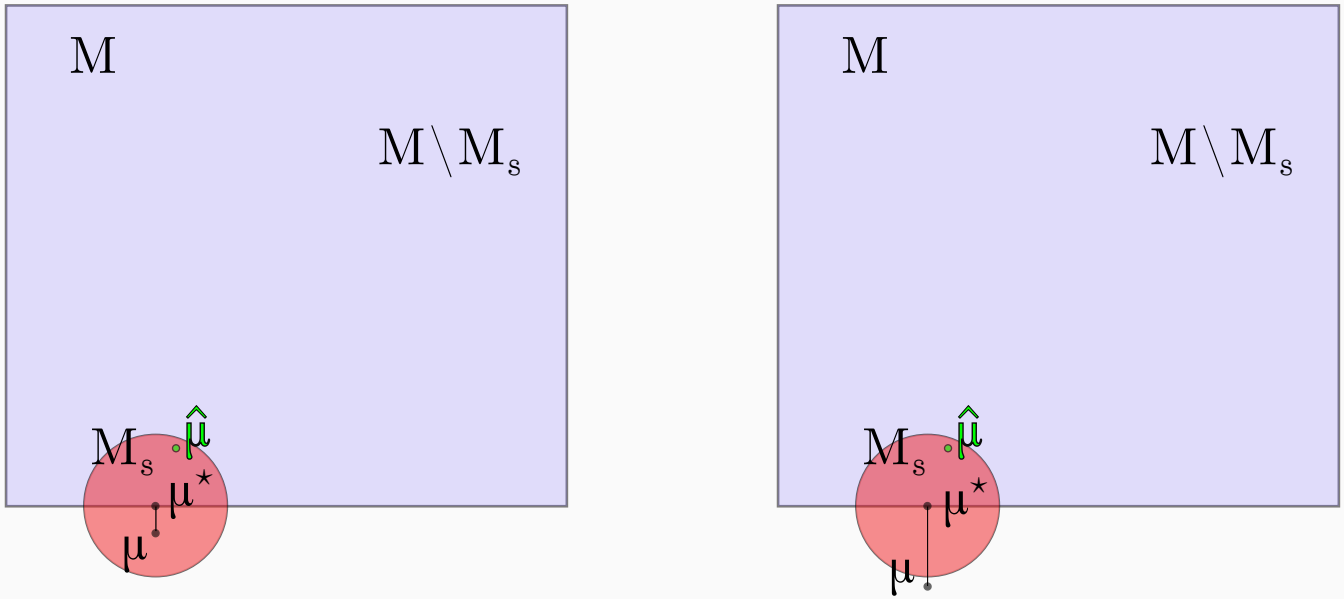

- Our error in estimating μ is bounded by a sum of two terms.
  - The critical radius s, i.e., the one satisfying $s^2/2\sigma \geq \mathrm{w}(\mathcal{M}_s^\circ) + s\sqrt{\frac{2\Sigma_n}{\delta n}}$.
  - The distance from μ to its best approximation in the model. Or really 3 times that.

  We showed this in the model selection lab using the Cauchy-Schwarz inequality.
- In convex models, we can say more.
  Our error in estimating $\mu^\star$ does not depend on its distance to μ.

Review

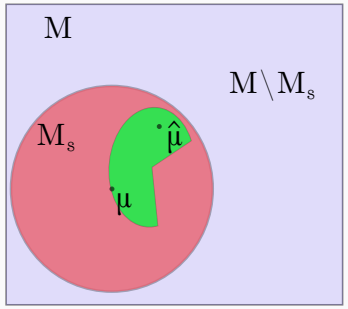

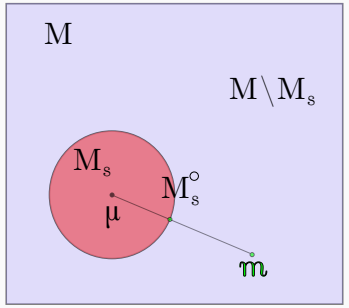

squared error loss

$$\hat{\mu} \quad \text{minimizes} \quad \ell(m) = \frac{1}{n}\sum_{i=1}^{n}\{\,Y_i - \mu(X_i)\}^2$$

among curves m in a convex set $\mathcal{M}$.

- If μ is in the model, that tells us it's one of the curves with loss as small as μ's.

$$\text{i.e.} \quad m = \hat{\mu} \quad \text{satisfies} \quad \ell(m) \le \ell(\mu) \quad \text{if} \quad \mu \in \mathcal{M}.$$

- To prove $\hat{\mu}$ is in the neighborhood $\mathcal{M}_s$, we show that ...
- ...none of these curves is in the neighborhood's complement $M \setminus \mathcal{M}_s$.

$$\hat{\mu} \in \mathcal{M}_s \quad \text{if} \quad \ell(m) > \ell(\mu) \quad \text{for all} \quad m \in \mathcal{M} \setminus \mathcal{M}_s.$$

- i.e. we show the *loss difference* is strictly positive for curves in the complement.
- That's true if it's positive for curves on the neighborhood's boundary $\mathcal{M}_s^\circ$.

$$\ell(m) - \ell(\mu) > 0 \quad \text{for all} \quad m \in \mathcal{M} \setminus \mathcal{M}_s \quad \text{if} \quad \ell(m) > \ell(\mu) \quad \text{for all} \quad m \in \mathcal{M}_s^\circ.$$

- And that boils down to the neighborhood's *squared radius* exceeding ...
- ...twice its boundary's *maximal inner product with noise* $\varepsilon = Y - m$.

$$\ell(m) - \ell(\mu) = s^2 - \langle \underbrace{Y - \mu}_{\varepsilon},\ m - \mu\rangle \ge s^2 - 2\max_{m \in \mathcal{M}_s^\circ}\langle Y - \mu,\ m - \mu\rangle \quad \text{for all} \quad m \in \mathcal{M}_s^\circ$$

- Then we do a little probability and get our error bound.

The Argument with no if

For any $\mu^\star \in \mathcal{M}$, we can expand our mean squared error difference as before.

$$\ell(m) - \ell(\mu^\star) = \|m - \mu^\star\|^2_{L_2(\mathrm{P_n})} - \frac{2}{n}\sum_{i=1}^{n} \varepsilon_i^\star \{m(X_i) - \mu^\star(X_i)\} \quad \text{for} \quad \varepsilon_i^\star = Y_i - \mu^\star(X_i).$$

But our new 'noise' $\varepsilon_i^\star$ doesn't have mean zero. It's our old noise ε_i, minus something.

$$\varepsilon_i^\star = \{Y_i - \mu(X_i)\} - \{\mu^\star(X_i) - \mu(X_i)\}.$$
$$\varepsilon_i \qquad \text{something}$$

So we can think of our mean squared error difference as having three terms:

$$\ell(m) - \ell(\mu^\star) = \|m - \mu^\star\|^2_{L_2(\mathrm{P_n})}$$ squared distance, like before;

$$- \frac{2}{n}\sum_{i=1}^{n} \varepsilon_i \{m(X_i) - \mu^\star(X_i)\}$$ a mean zero term, like before;

$$+ \frac{2}{n}\sum_{i=1}^{n} \{\mu^\star(X_i) - \mu(X_i)\}\{m(X_i) - \mu^\star(X_i)\}$$ and something else.

We can use our argument, ignoring the new term, if that term is always *non-negative*.

Why?

Why.

$$\ell(m) - \ell(\mu^\star) = \|m - \mu^\star\|^2_{L_2(\mathrm{Pn})} - \frac{2}{n}\sum_{i=1}^{n} \varepsilon_i \{m(X_i) - \mu^\star(X_i)\} + \frac{2}{n}\sum_{i=1}^{n}\{\mu^\star(X_i) - \mu(X_i)\}\{m(X_i) - \mu^\star(X_i)\}$$

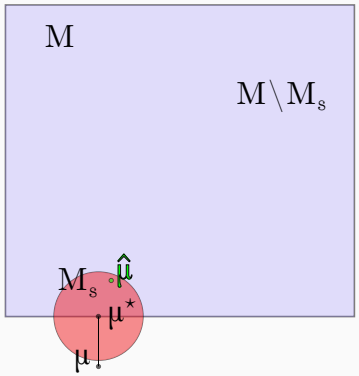

We want to show that if distance from m to $\mu^\star$ is big enough, it wins.

- In particular, it wins in the sense that the loss difference $\ell(m) - \ell(\mu^\star)$ is positive.
- That implies distance from $\hat{\mu}$ to $\mu^\star$ is smaller, as distance doesn't win in that case.

If this new term is non-negative, it helps distance win.

- If the loss difference is positive when we ignore a non-negative term ...
- ...then it's still positive when we don't.

$$\ell(m) - \ell(\mu^\star) > 0 \quad \text{if} \quad \|m - \mu^\star\|^2_{L_2(\mathrm{Pn})} - \frac{2}{n}\sum_{i=1}^{n} \varepsilon_i \{m(X_i) - \mu^\star(X_i)\} > 0 \qquad \text{what we're used to}$$

$$\text{and} \quad \frac{2}{n}\sum_{i=1}^{n}\{\mu^\star(X_i) - \mu(X_i)\}\{m(X_i) - \mu^\star(X_i)\} \geq 0 \qquad \text{new term}$$

This only works if the new term is non-negative. Can we choose $\mu^\star \in \mathcal{M}$ so it is?

We can

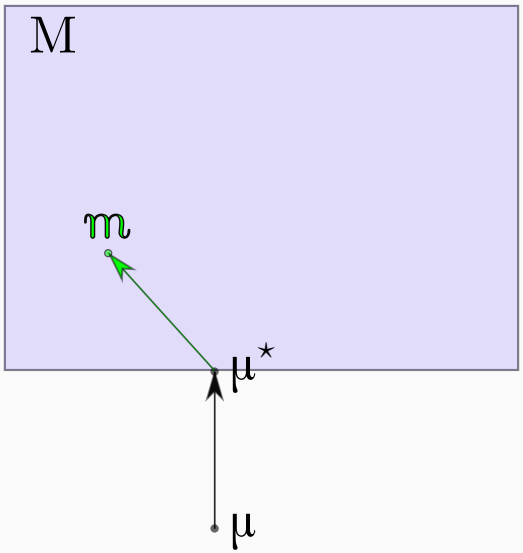

The new term is always non-negative when we compare to the *best approximation* to μ in the model,

$$\mu^\star = \underset{m \in \mathcal{M}}{\operatorname{argmin}} \|m - \mu\|^2_{L_2(\mathrm{P_n})} \quad \text{satisfies} \quad \frac{2}{n}\sum_{i=1}^{n}\{\mu^\star(X_i) - \mu(X_i)\}\{m(X_i) - \mu^\star(X_i)\}$$

$$\text{or in vector notation} \quad \frac{2}{n}\langle \mu^\star - \mu, m - \mu^\star \rangle_2 \geq 0 \quad \text{for all} \quad m \in \mathcal{M}.$$

It's proportional to the dot product between two vectors: $\mu \to \mu^\star$ and $\mu^\star \to m$.

- When the model $\mathcal{M}$ is convex, these vectors are always in the same direction.
- They both point 'in' to the model. That means the dot product is non-negative.

Proof

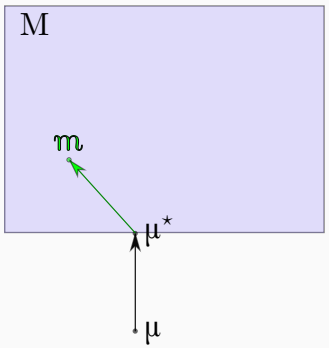

Claim. For any convex set $\mathcal{M}$ in an inner product space, [1]

$$\mu^\star = \operatorname*{argmin}_{m \in \mathcal{M}} \|m - \mu\| \quad \text{satisfies}$$

$$\langle \mu^\star - \mu,\ m - \mu^\star \rangle \geq 0 \quad \text{for all curves} \quad m \in \mathcal{M}.$$

Proof. Let $m_\lambda = \lambda(m - \mu^\star) + \mu^\star$.

$$\begin{aligned} \|m_\lambda - \mu\|^2 &= \langle \lambda(m - \mu^\star) + (\mu^\star - \mu),\ \lambda(m - \mu^\star) + (\mu^\star - \mu) \rangle \\ &= \lambda^2 \|m - \mu^\star\|^2 + \|\mu^\star - \mu\|^2 + 2\lambda \langle m - \mu^\star,\ \mu^\star - \mu \rangle. \end{aligned}$$

Because $m_\lambda \in \mathcal{M}$, it follows that this is at least as large as $\|\mu - \mu^\star\|^2$, so

$$0 \leq \lambda^2 \|m - \mu^\star\|^2 + 2\lambda \langle m - \mu^\star,\ \mu^\star - \mu \rangle$$

and therefore, dividing by $\lambda > 0$, that

$$0 \leq \lambda \|m - \mu^\star\|^2 + 2 \langle m - \mu^\star,\ \mu^\star - \mu \rangle.$$

Because this holds for arbitrarily small $\lambda > 0$, it must also hold for $\lambda = 0$.

[1] An inner product space is a vector space with a norm $\|u\| = \sqrt{\langle u, u \rangle}$ induced by an inner product $\langle u, v \rangle$.

That's not true for other choices

When $\mu^\star \in \mathcal{M}$ isn't the closest point to μ,
these vectors can point in opposite directions.
That is, this dot product can be negative for some $m \in \mathcal{M}$.

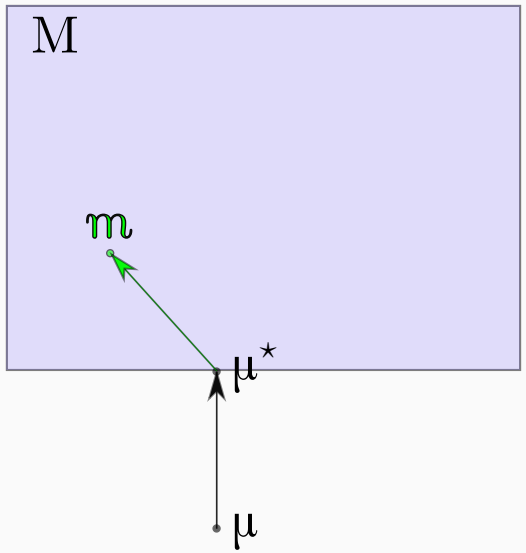

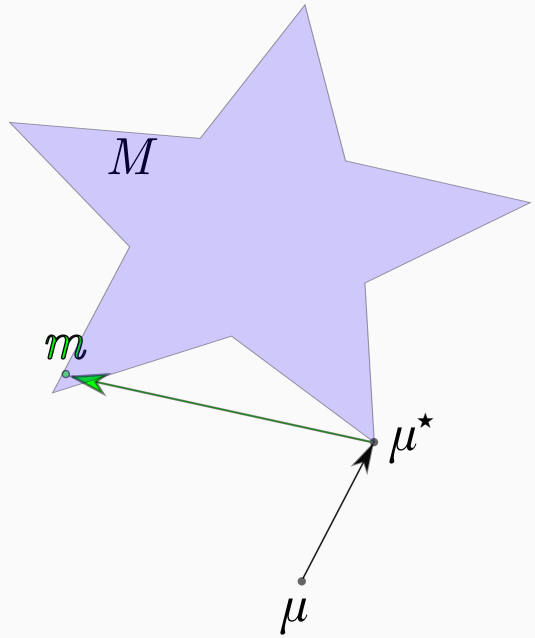

The same thing can happen *for the closest point* in a non-convex model.

Summary

When we use a convex model, the least squares estimator $\hat{\mu}$ converges to the model's closest point to μ. This generalizes our result without misspecification.

- If μ is in the model, that closest point is μ.
- Otherwise, it's something else.

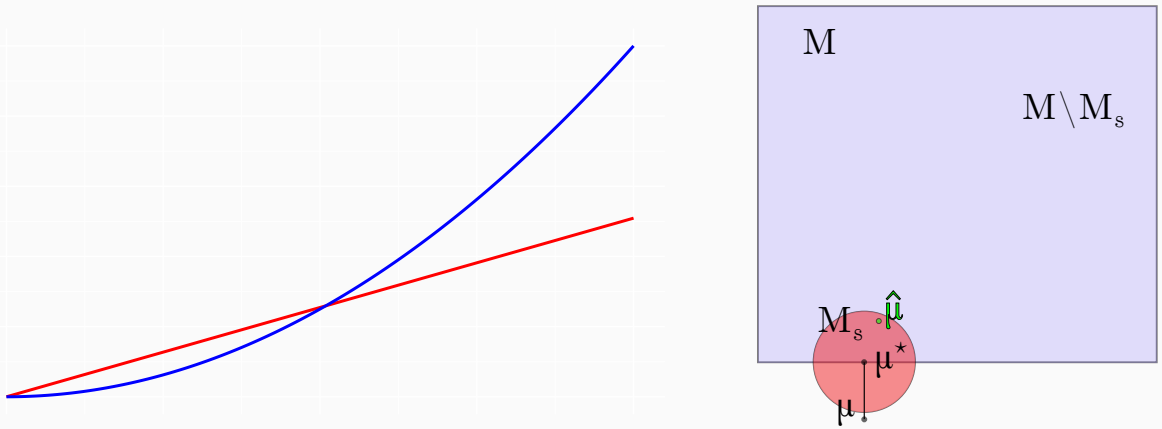

We can bound our estimator's distance to that closest point $\mu^\star$ just like we've been bounding distance to μ when we assumed it was in the model.

$$\|\hat{\mu} - \mu^\star\|_{L_2(\mathrm{P_n})} < s + 2\sqrt{\frac{2\Sigma_n}{\delta n}} \text{ w.p. } 1-\delta \text{ if } s^2/2\sigma \geq \mathrm{w}(\mathcal{M}_s)$$

for $\mathcal{M}_s = \{m \in \mathcal{M} : \|m - \mu^\star\|_{L_2(\mathrm{P_n})} \leq s\}$ and $\Sigma_n = \sigma\{1 + 2\log(2n)\}$

if $Y_i = \mu(X_i) + \varepsilon_i$ for $\varepsilon_i \overset{iid}{\sim} N(0, \sigma^2)$ **for some function μ.**

Misspecification

Examples

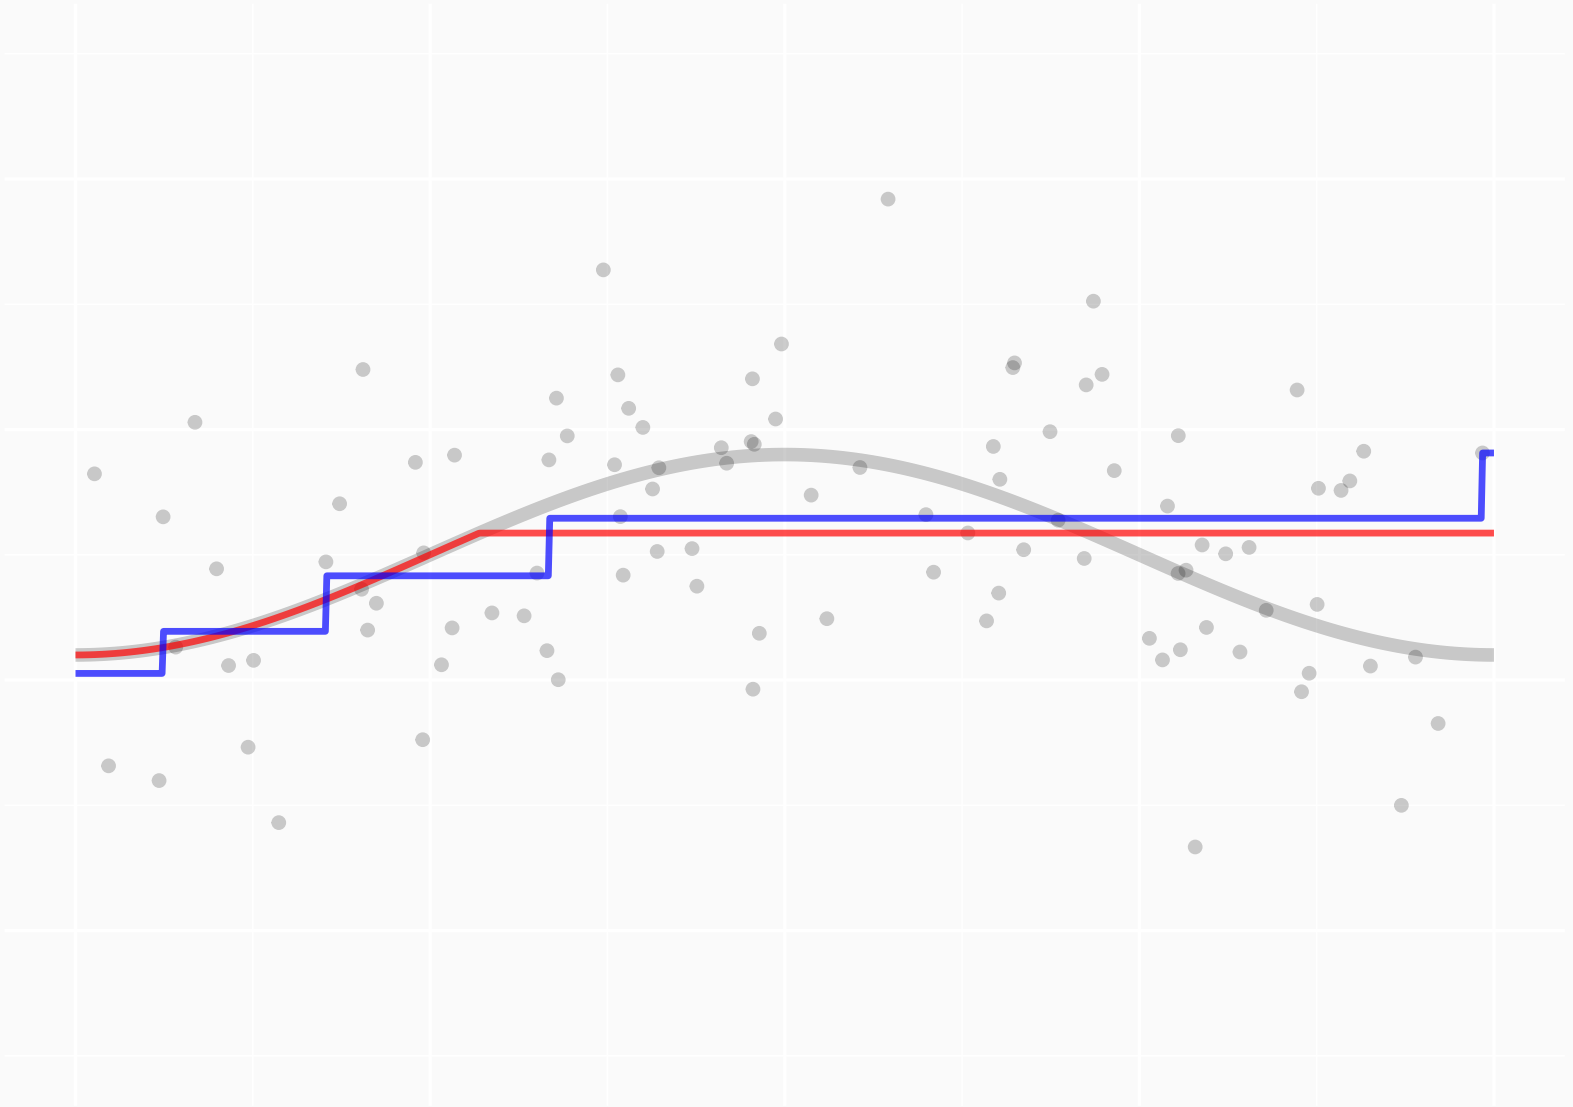

Figure 1: Increasing Curves ($n = 100$.)

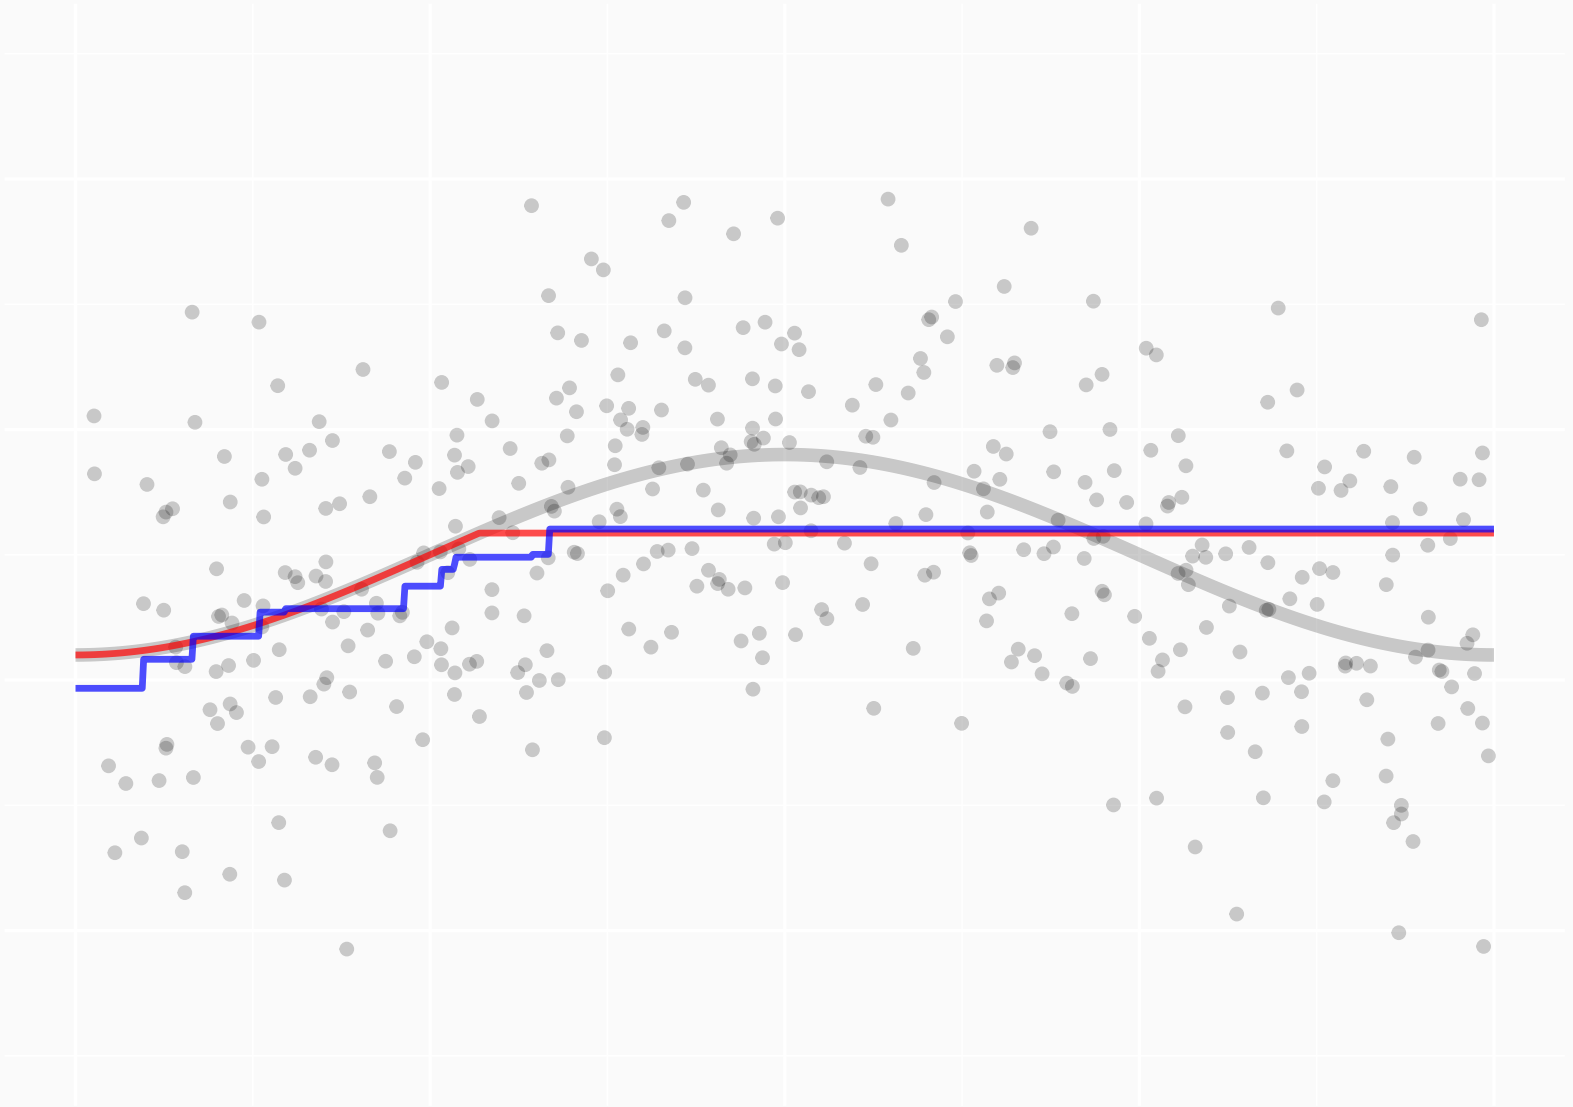

Figure 2: Increasing Curves ($n = 400$.)

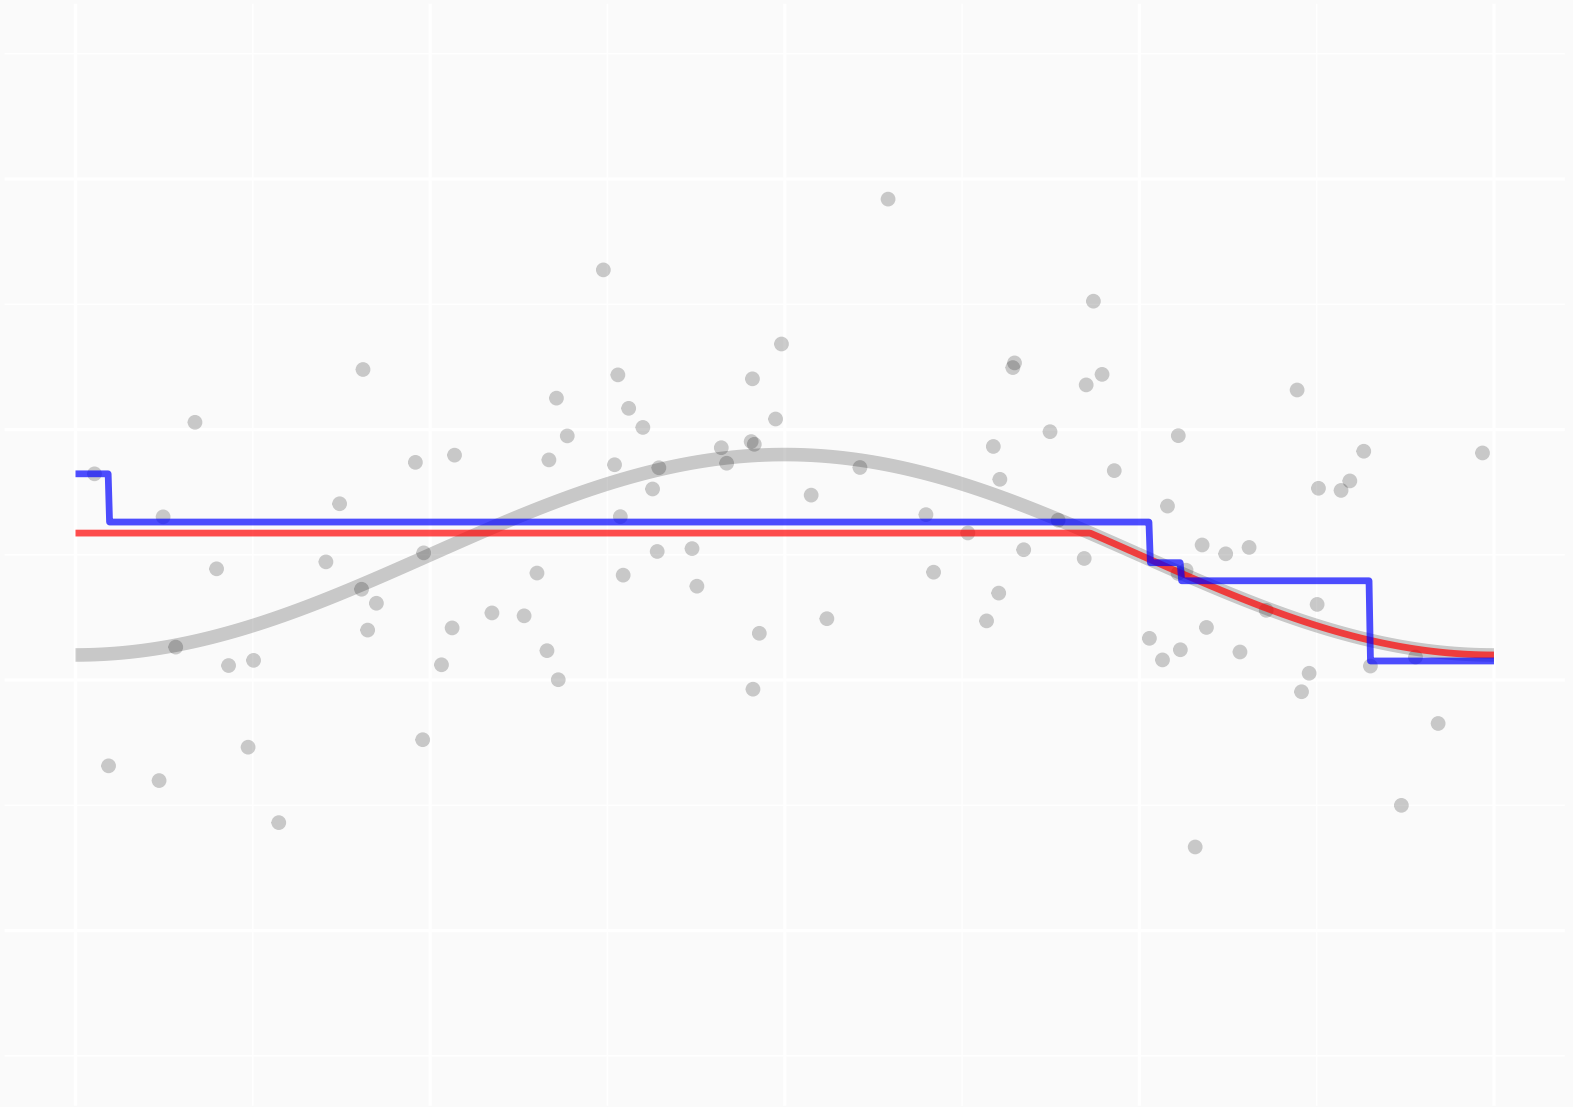

Figure 3: Decreasing Curves ($n = 100$.)

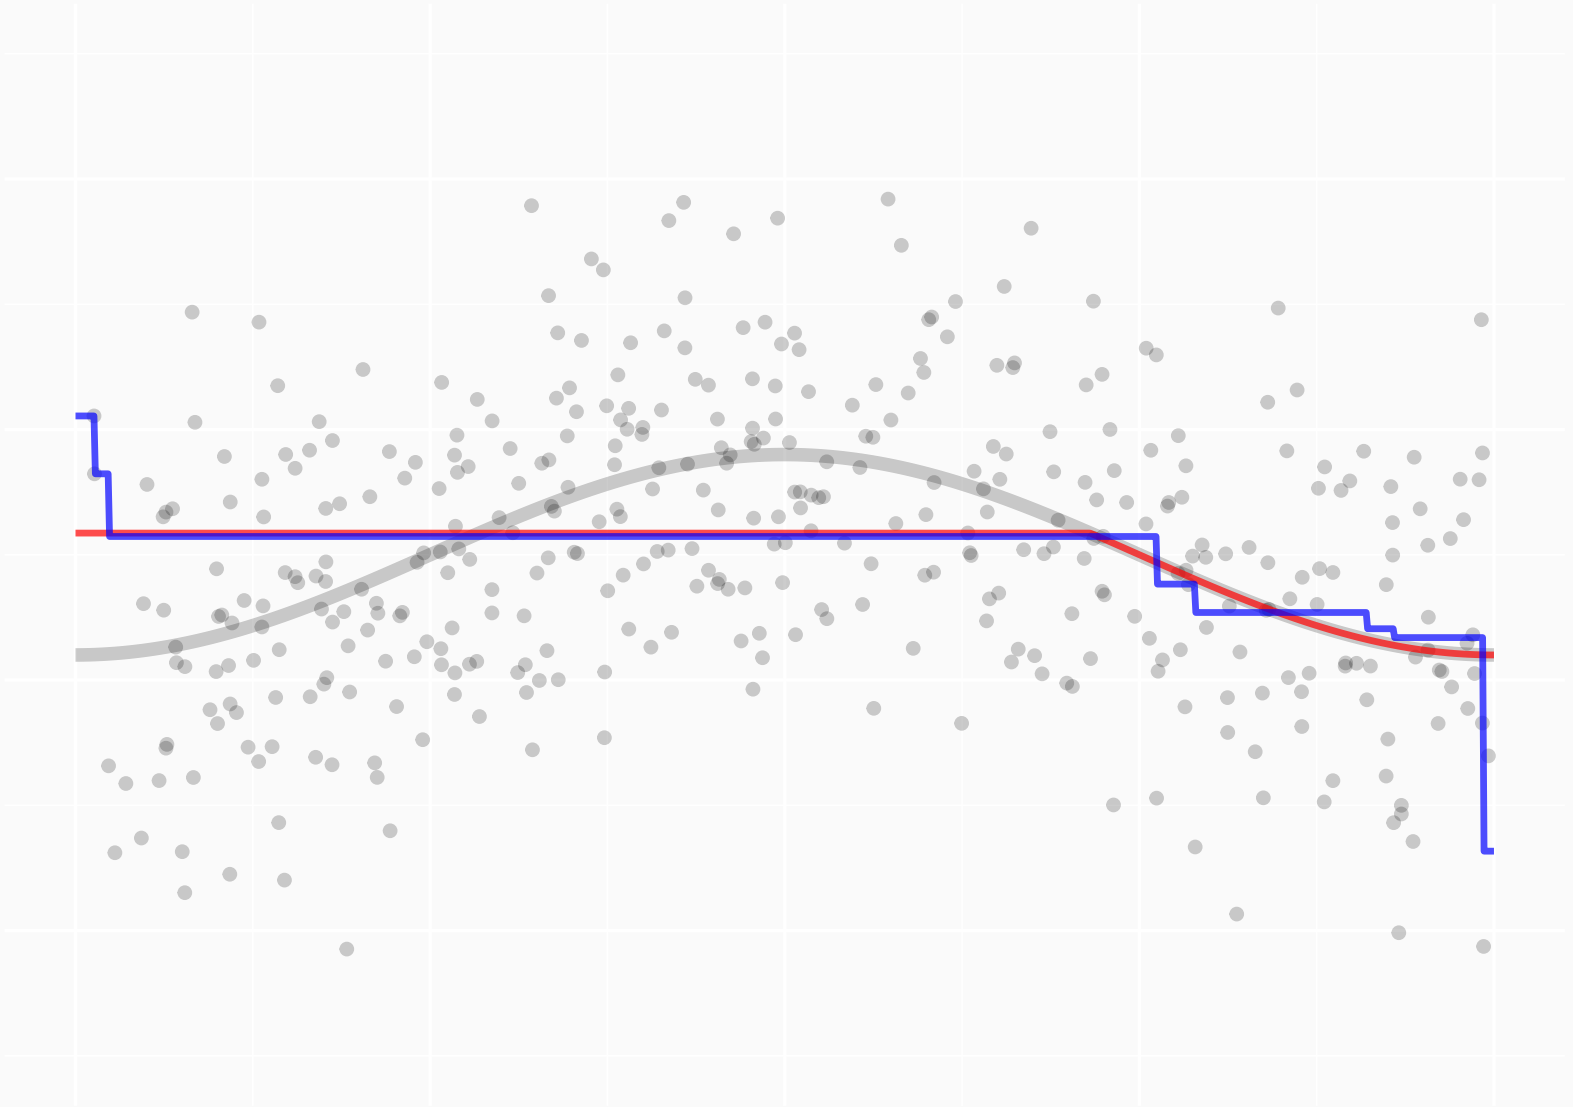

Figure 4: Decreasing Curves ($n = 400$.)

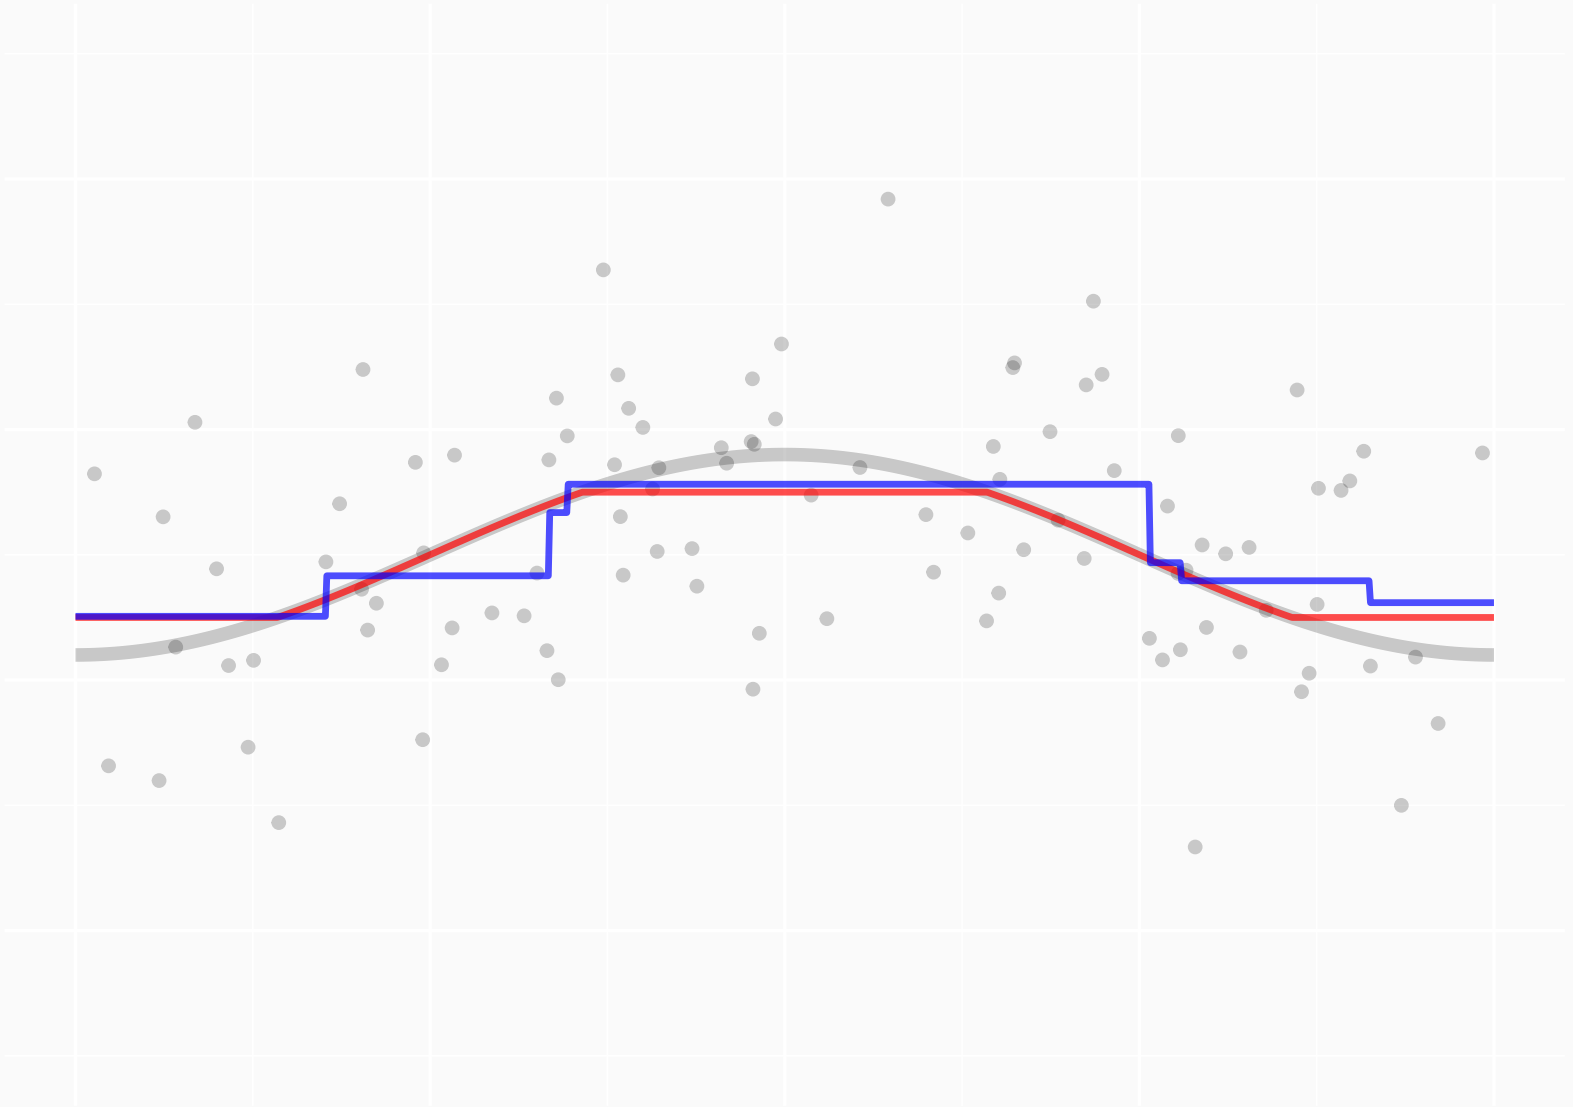

Figure 5: Bounded Variation Curves: $\rho_{\mathsf{TV}} \leq 1$ $(n = 100.)$

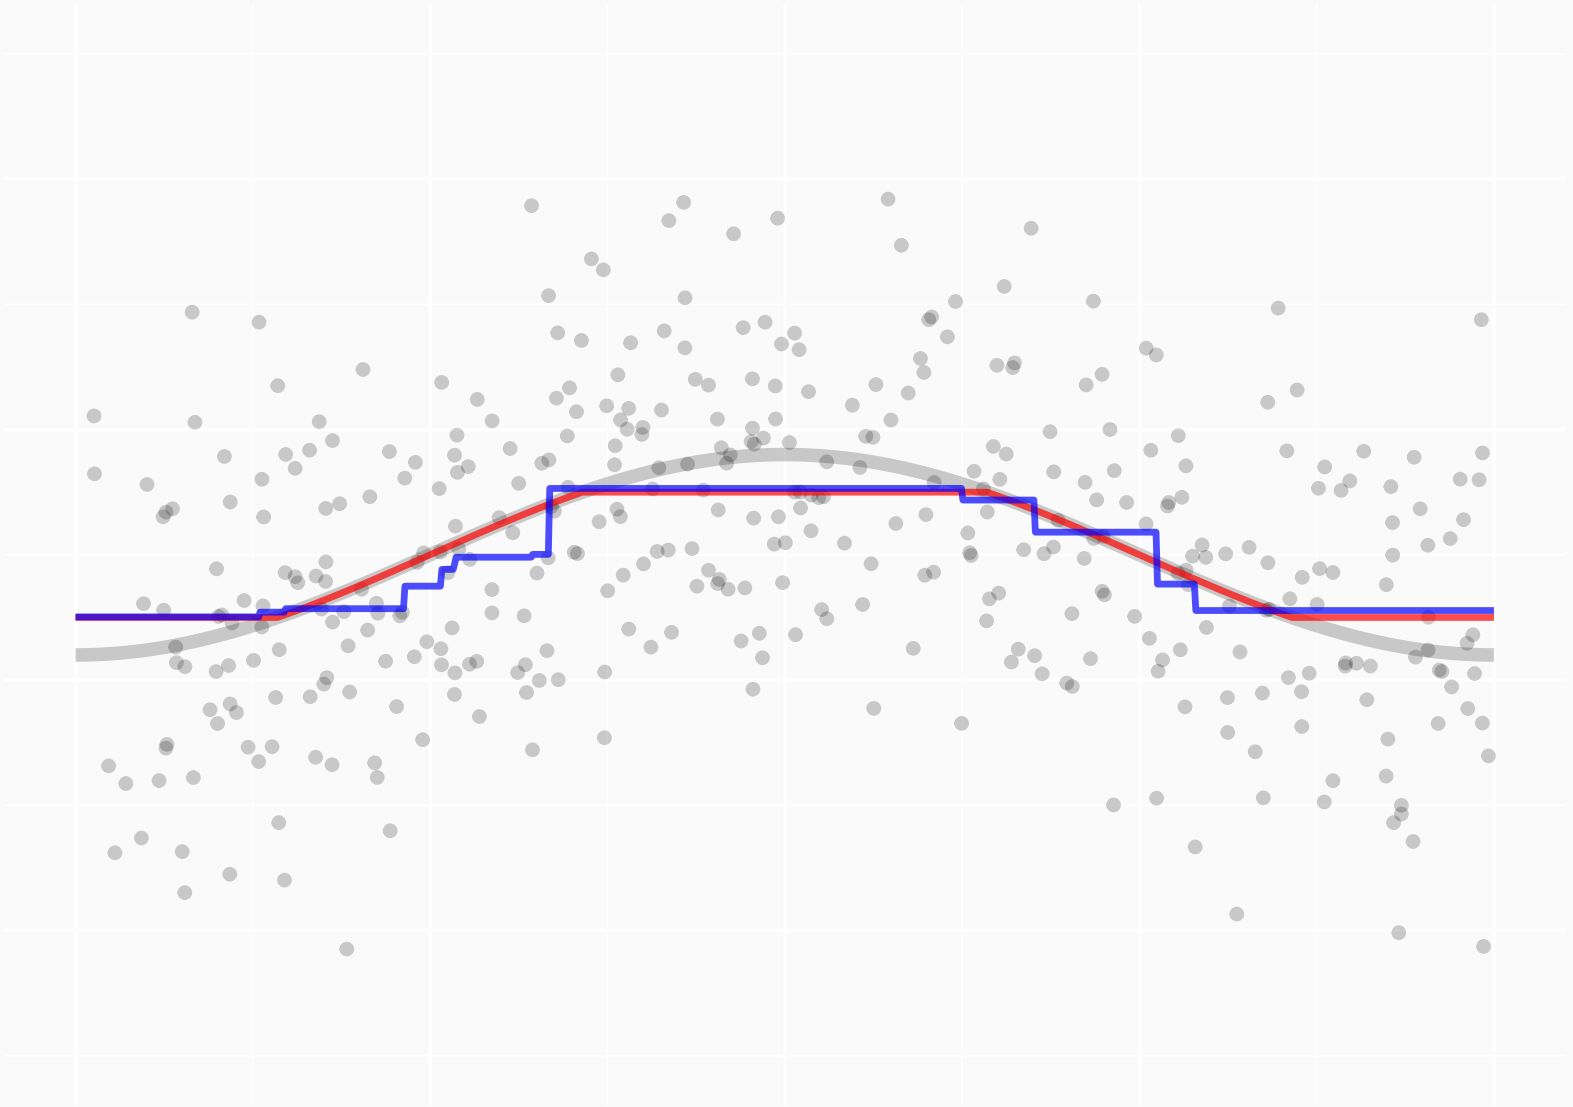

Figure 6: Bounded Variation Curves: $\rho_{\mathsf{TV}} \leq 1$. ($n = 400$.)

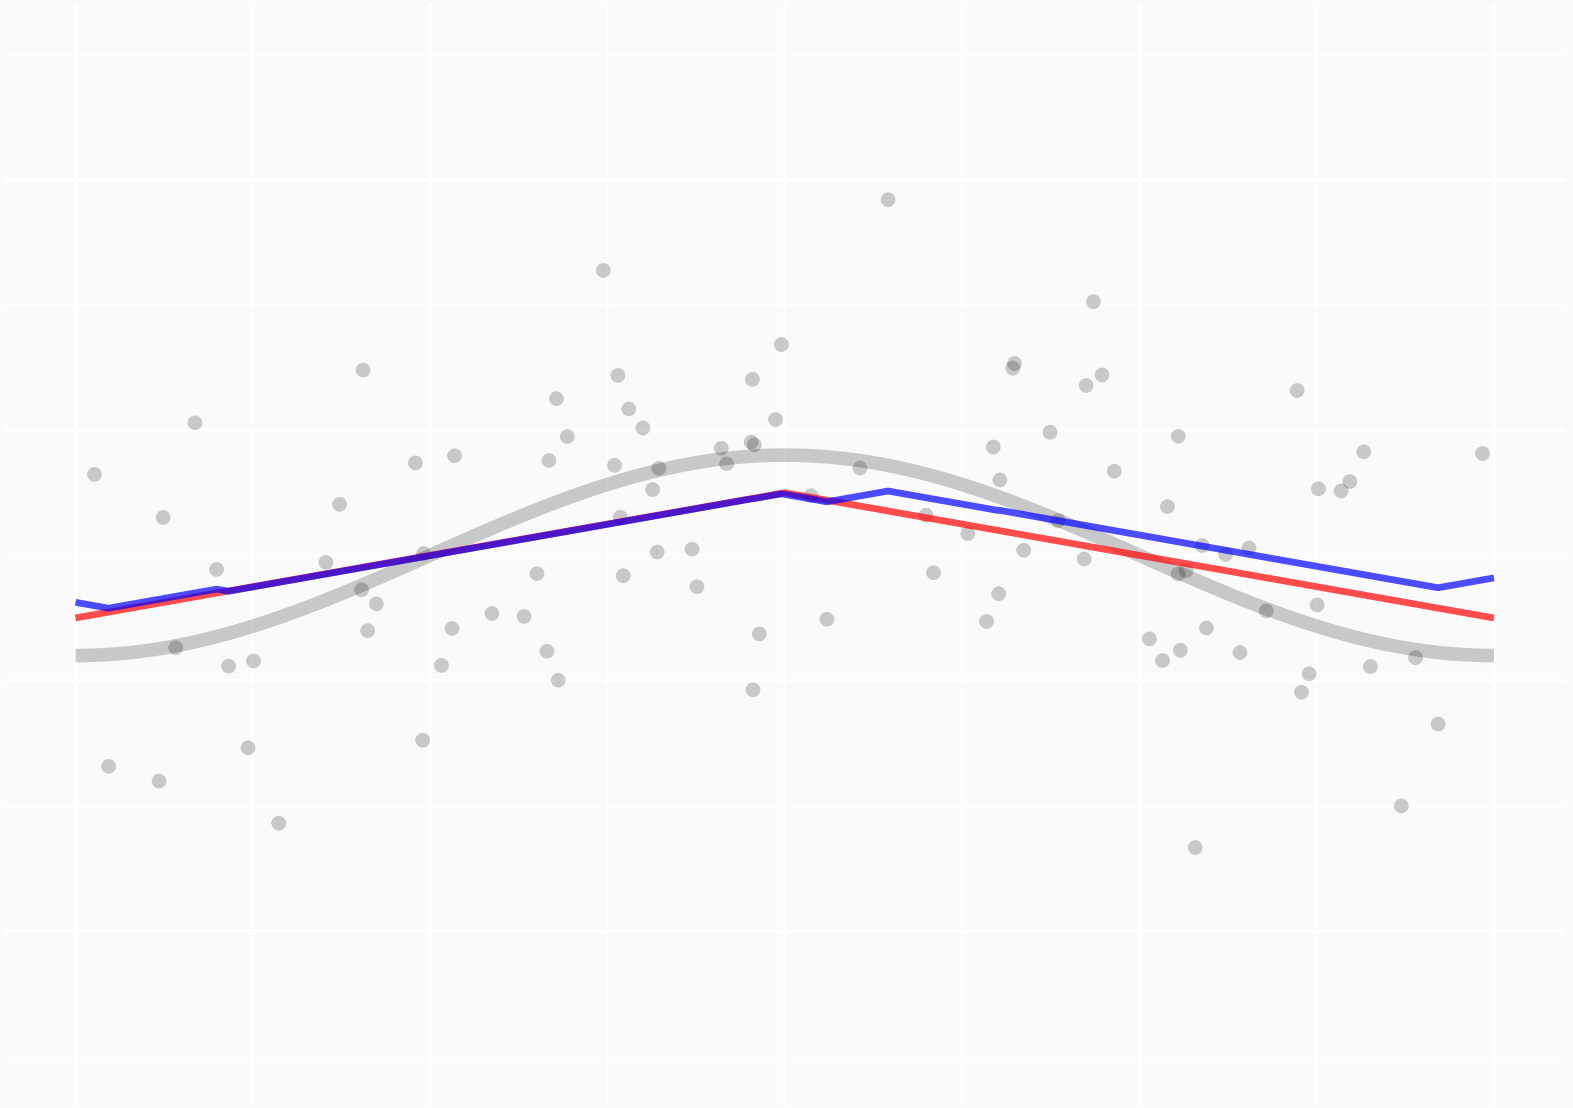

Figure 7: Lipschitz Curves: $\rho_{\text{Lip}} \leq 1$ $(n = 100.)$

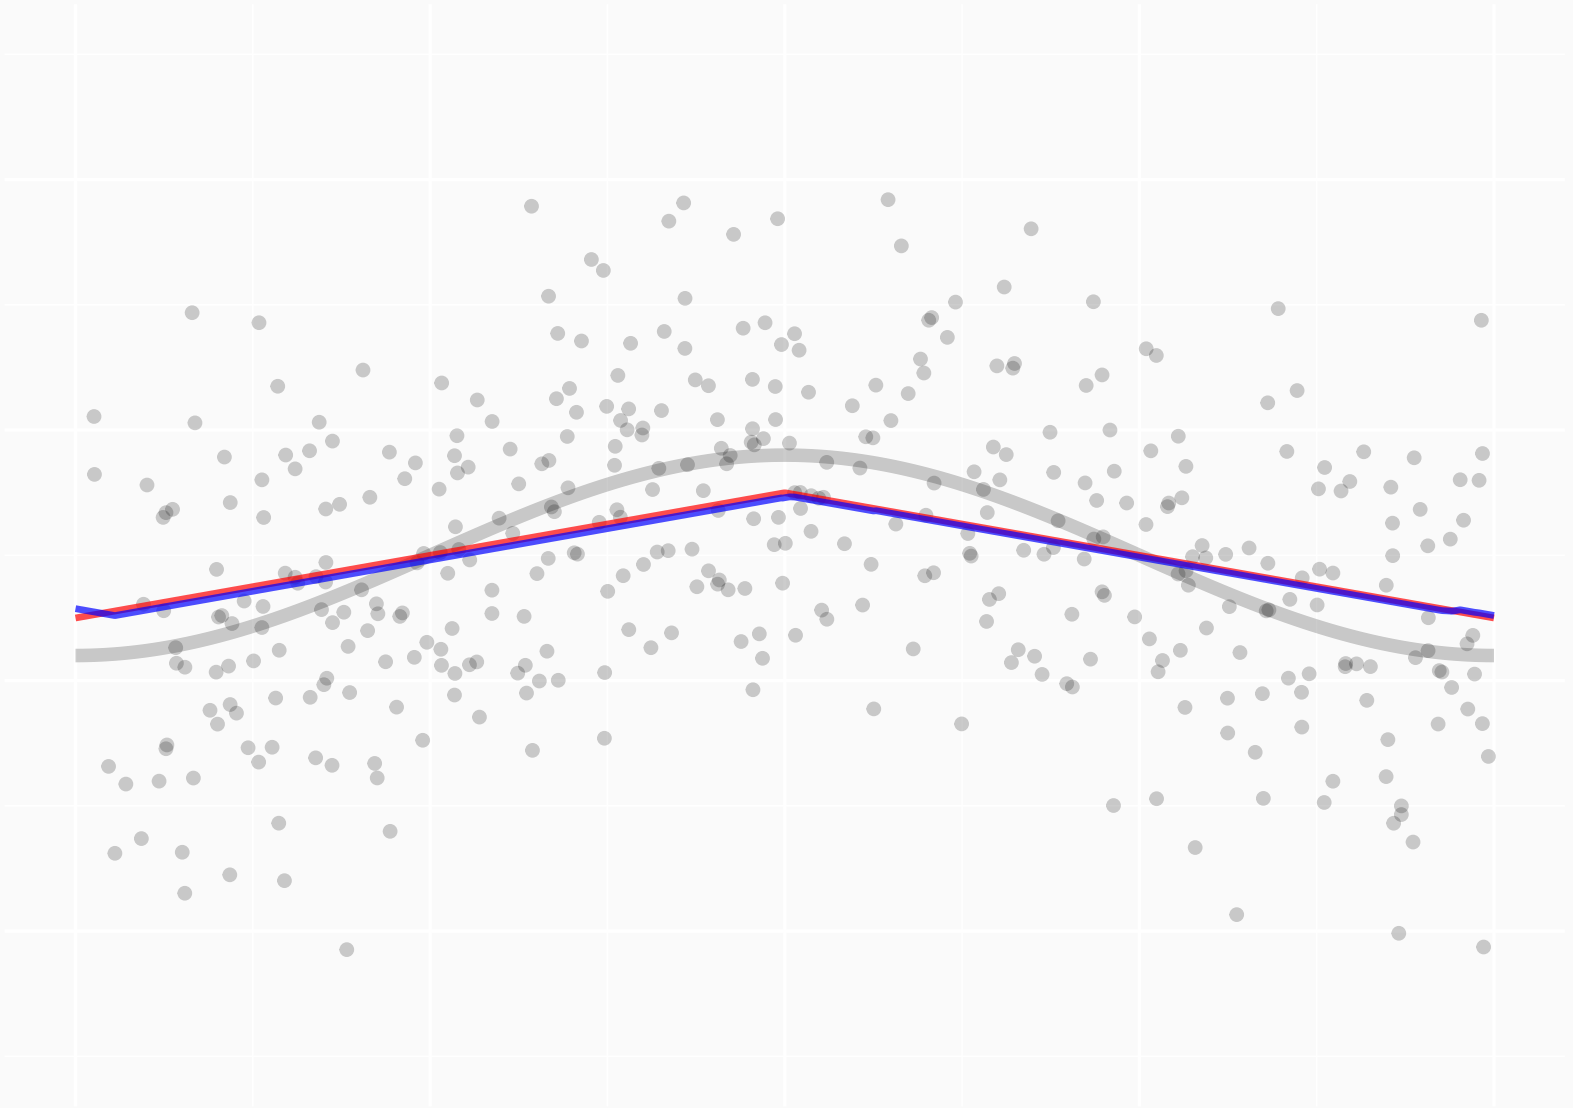

Figure 8: Lipschitz Curves: $\rho_{\text{Lip}} \leq 1$ $(n = 400.)$

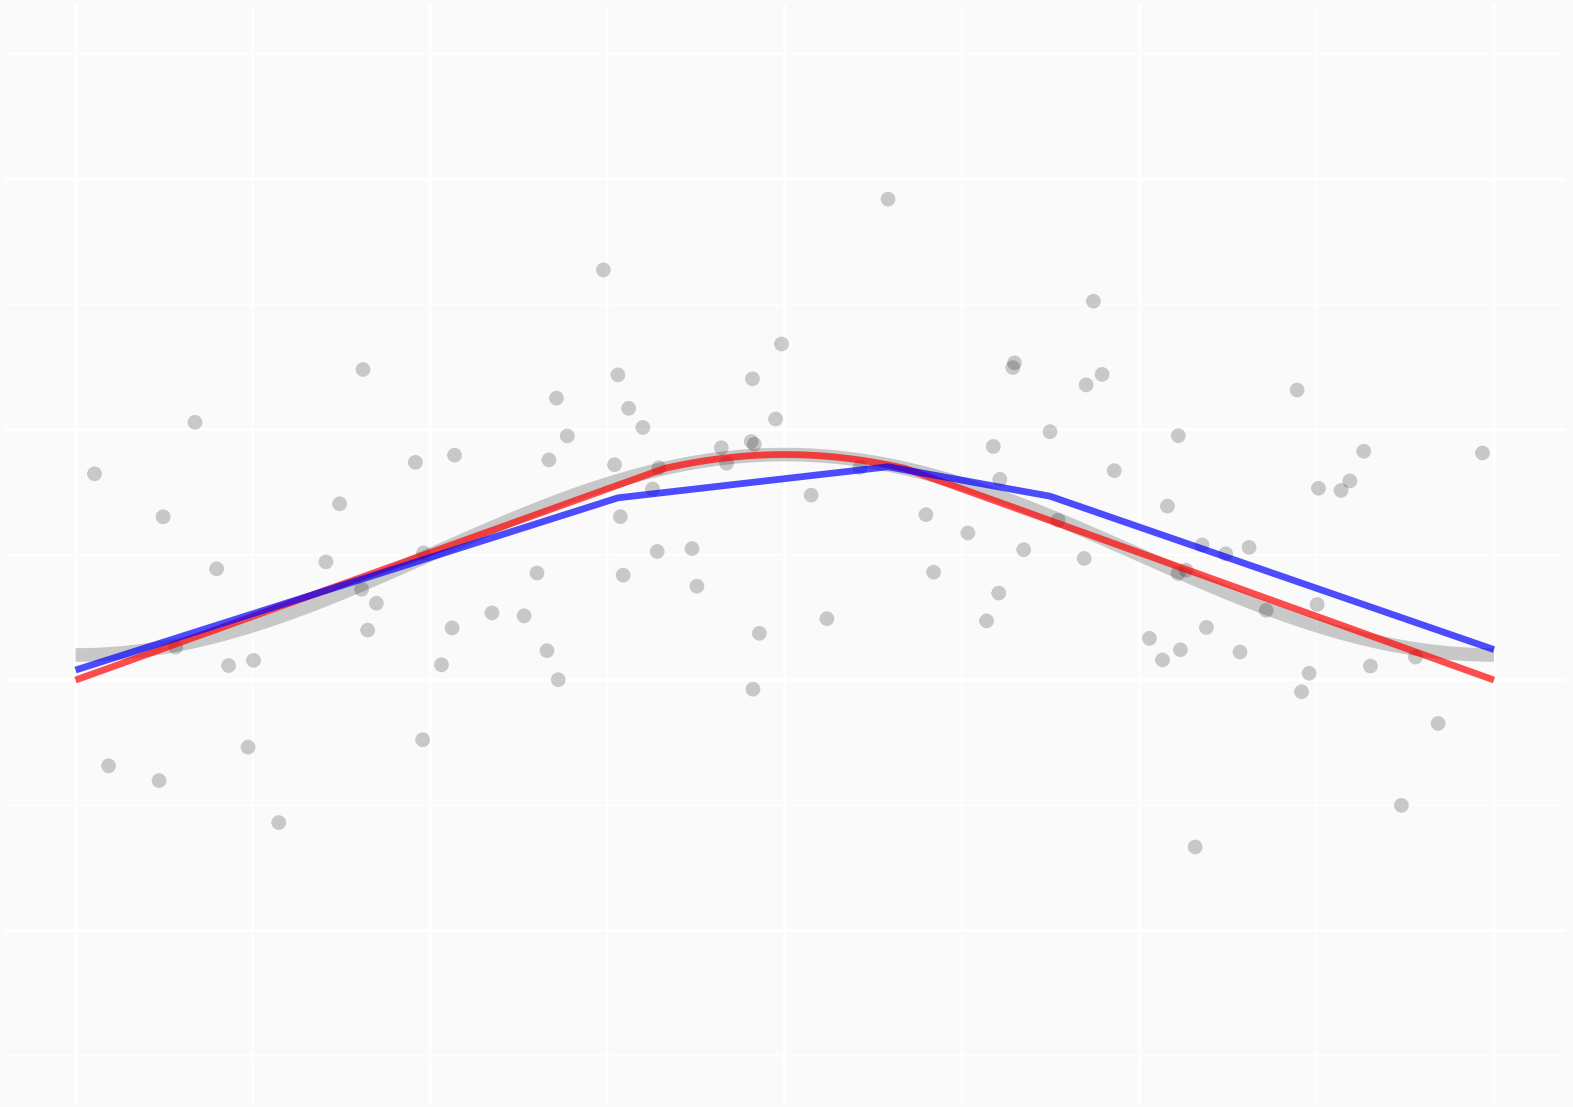

Figure 9: Concave Curves ($n = 100$.)

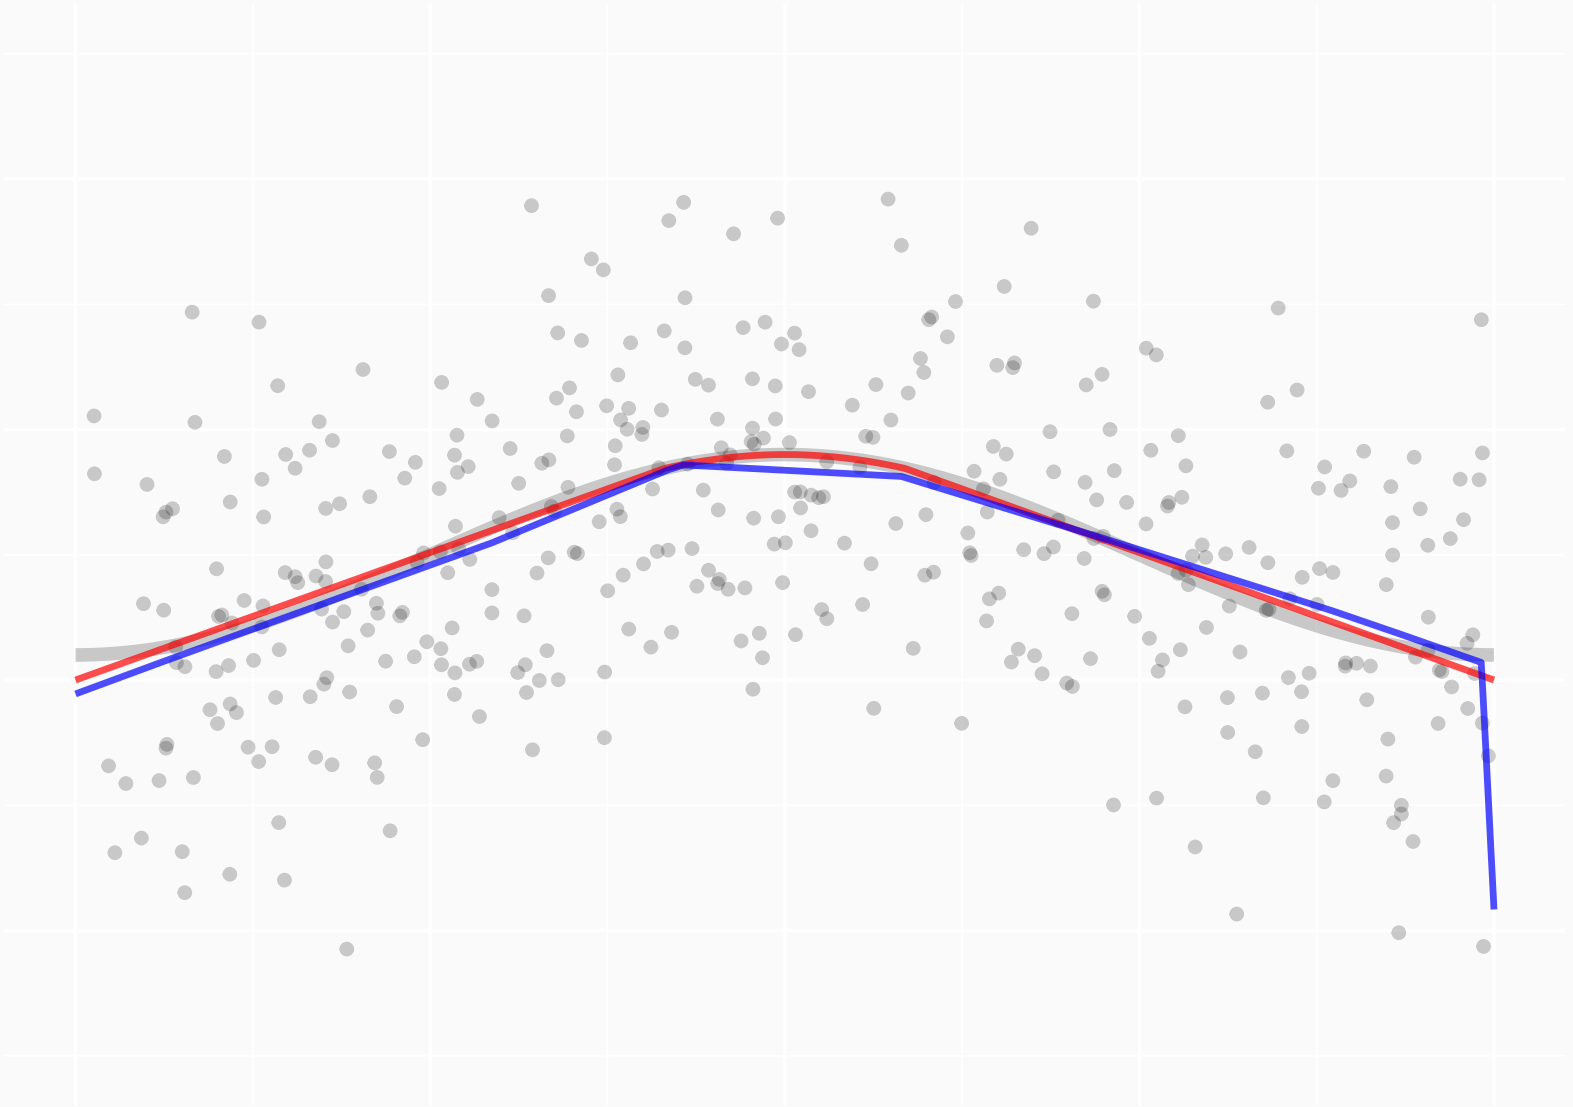

Figure 10: Concave Curves ($n = 400$.)

Non-Gaussian Noise

Starting Point

$$\ell(m) - \ell(\mu^\star) = \|m - \mu^\star\|^2_{L_2(\mathrm{P_n})} \qquad \text{squared distance}$$

$$- \frac{2}{n}\sum_{i=1}^{n} \varepsilon_i \{m(X_i) - \mu^\star(X_i)\} \qquad \text{a mean zero term}$$

$$+ \frac{2}{n}\sum_{i=1}^{n} \{\mu^\star(X_i) - \mu(X_i)\}\{m(X_i) - \mu^\star(X_i)\} \qquad \text{a non-negative term.}$$

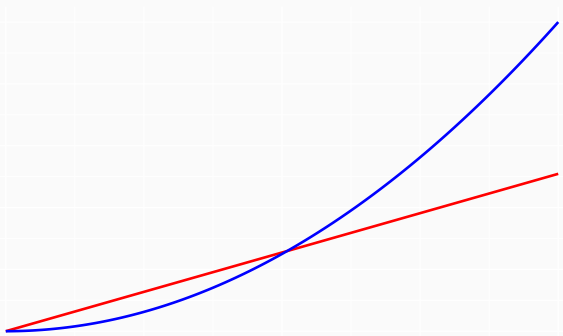

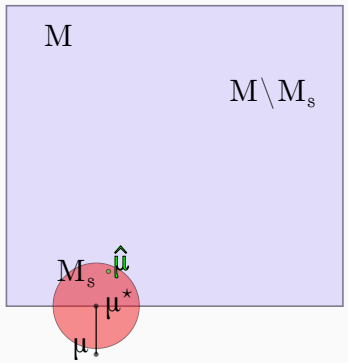

We can bound error using a corresponding *width*, no matter how noise is distributed.

$$\|\hat{\mu} - \mu^\star\|_{L_2(\mathrm{P_n})} < s + 2\sqrt{\frac{2\Sigma_n}{\delta n}} \quad \text{w.p. } 1-\delta \text{ for } \frac{s^2}{2} \geq \mathrm{w}_\epsilon(\mathcal{M}_s)$$

where $\mathrm{w}_\epsilon(\mathcal{V}) = \mathrm{E} \max_{v \in \mathcal{V}} \langle \epsilon, v \rangle_{L_2(\mathrm{P_n})}$ and $\Sigma_n = \mathrm{E} \max_{i \in 1 \ldots n} \varepsilon_i^2$.

This bound depends on the model $\mathcal{M}$ and the distribution of the noise ε in a complex, entangled way: through the width $\mathrm{w}_\varepsilon(\mathcal{M}_s)$.

Plan for Today

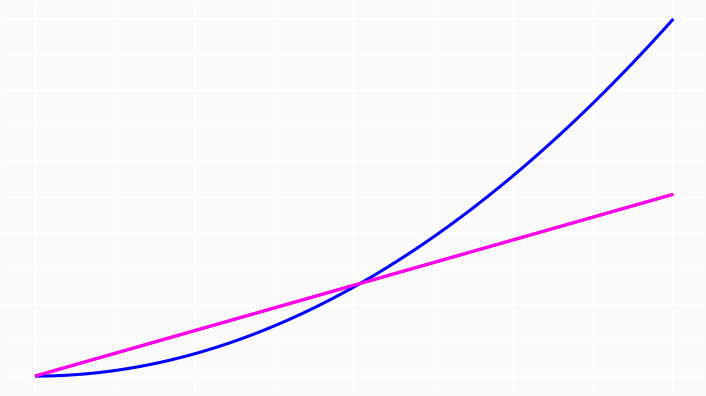

To disentangle the impact of the model and noise distribution, we'll bound this width in terms of gaussian width.

$$\mathrm{w}_\epsilon(\mathcal{M}_s) \leq \alpha \mathrm{w}(\mathcal{M}_s)$$

for α depending on ε but not $\mathcal{M}$ or s.

At the heart of this comparison $\mathrm{w}_\epsilon(\cdot) \leq \alpha \mathrm{w}(\cdot)$ are two ideas.

1. **Symmetrization**. We'll substitute for ϵ_i a variant that's symmetric around zero.

   $$\epsilon_i \to \epsilon_i - \epsilon_i' \quad \text{where} \quad \epsilon_i' \text{ is an independent copy of } \epsilon_i$$

   This substitution *increases* width: $\mathrm{w}_\epsilon(\cdot) \leq \mathrm{w}_{\epsilon - \epsilon'}(\cdot)$.
2. **Contraction**. We'll substitute a gaussian vector for our symmetrized noise $\epsilon - \epsilon'$. We can bound the impact of this substitution in a model-invariant way.

   $$\mathrm{w}_{\epsilon - \epsilon'}(\cdot) \leq \sqrt{2\pi} M_n \times \mathrm{w}(\cdot) \quad \text{for} \quad M_n = \mathrm{E} \max_{i \in 1 \ldots n} |\varepsilon_i|$$

This lets us re-use our gaussian width calculations to analyze regression with any noise distribution.

A Simple Consequence: Width Comparison implies Radius Comparison

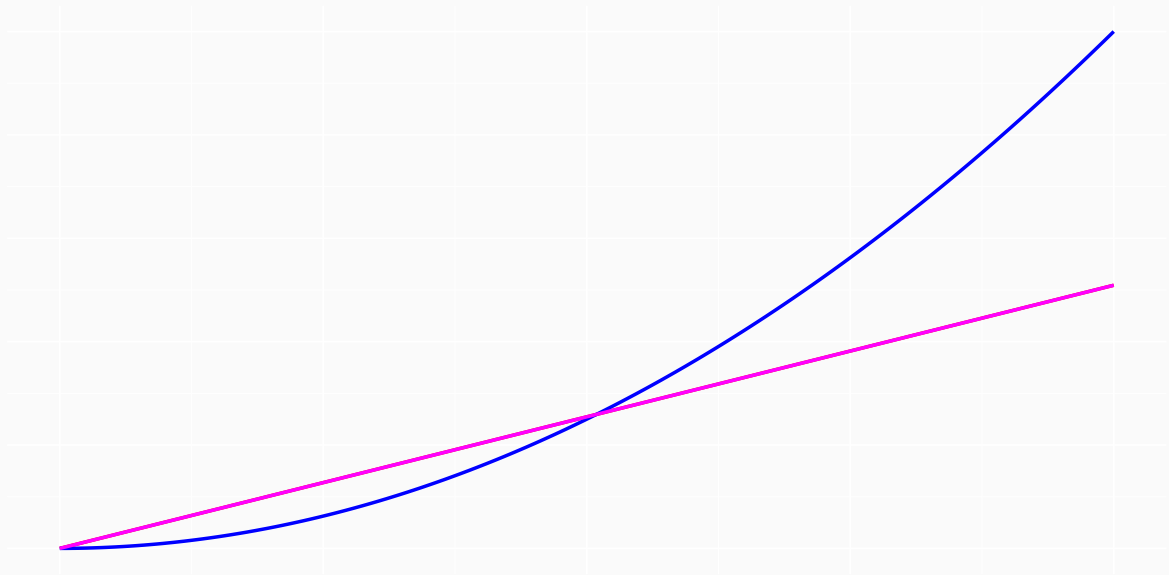

- If you have a width comparison $\mathrm{w}_\epsilon \leq \alpha \mathrm{w}_\eta$ for some $\alpha \geq 1$.
- This implies a radius comparison $s_\epsilon \leq \alpha s_\eta$ for all convex models $\mathcal{M}$.

$$s_\epsilon = \alpha s_\nu \quad \text{satisfies} \quad \frac{s_\epsilon^2}{2} \geq \mathrm{w}_\varepsilon(\mathcal{M}_{s_\epsilon}) \quad \text{if} \quad \frac{s_\eta^2}{2} \geq \mathrm{w}_\eta(\mathcal{M}_{s_\eta}) \quad \text{for} \quad \text{convex } \mathcal{M}$$
$$\text{and} \quad \mathrm{w}_\varepsilon \leq \alpha \mathrm{w}_\eta \quad \text{for} \quad \alpha \geq 1.$$

- *Interpretation.*
  The noise ε makes regression at most 'α times harder' than the noise η.
- *This is simplistic and 'lossy'.*
  For most models, our width comparison implies a better radius comparison.

Proof: Width Comparisons imply Radius Comparisons

Claim. If $\mathrm{w}_\varepsilon \leq \alpha\, \mathrm{w}_\eta$ for $\alpha \geq 1$, then for any convex model $\mathcal{M}$, the critical radius using noise ε is at most α times the critical radius using noise η, i.e.

$$\frac{(\alpha s)^2}{2} \geq \mathrm{w}_\varepsilon(\mathcal{M}_{\alpha s}) \quad \text{if} \quad \frac{s^2}{2} \geq \mathrm{w}_\eta(\mathcal{M}_s) \quad \text{and} \quad \mathrm{w}_\varepsilon \leq \alpha\, \mathrm{w}_\eta \quad \text{for} \quad \alpha \geq 1.$$

Proof. If $s^2/2 \geq \mathrm{w}_\eta(\mathcal{M}_s)$, then

Proof: Width Comparisons imply Radius Comparisons

Claim. If $\mathrm{w}_\varepsilon \le \alpha\, \mathrm{w}_\eta$ for $\alpha \ge 1$, then for any convex model $\mathcal{M}$, the critical radius using noise ε is at most α times the critical radius using noise η, i.e.

$$\frac{(\alpha s)^2}{2} \ge \mathrm{w}_\varepsilon(\mathcal{M}_{\alpha s}) \quad \text{if} \quad \frac{s^2}{2} \ge \mathrm{w}_\eta(\mathcal{M}_s) \quad \text{and} \quad \mathrm{w}_\varepsilon \le \alpha\, \mathrm{w}_\eta \quad \text{for} \quad \alpha \ge 1.$$

Proof. If $s^2/2 \ge \mathrm{w}_\eta(\mathcal{M}_s)$, then

$$\begin{aligned} \alpha s/2 &\ge \alpha\, \mathrm{w}_\eta(\mathcal{M}_s)/s && \text{multiplying both sides by } \alpha/s \\ &\ge \alpha\, \mathrm{w}_\eta(\mathcal{M}_{\alpha s})/(\alpha s) && \text{using sublinearity of } f(s) = \mathrm{w}_\eta(\mathcal{M}_s) \\ &\ge \mathrm{w}_\varepsilon(\mathcal{M}_{\alpha s})/(\alpha s) && \text{using our premise } \alpha\, \mathrm{w}_\eta \ge \mathrm{w}_\varepsilon. \end{aligned}$$

Multiplying both sides by αs, we get our claim.

Summary

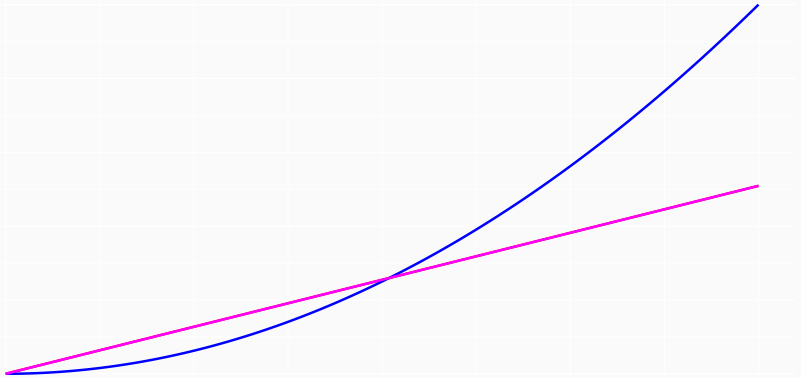

Where we are. We have a bound that depends on the model $\mathcal{M}$ and the distribution of the noise ε in a complex and entangled way.

$$\|\hat{\mu} - \mu^\star\|_{L_2(\mathrm{P_n})} < s_\epsilon + 2\sqrt{\frac{2\Sigma_n}{\delta n}} \quad \text{w.p. } 1-\delta \text{ for } \quad \frac{s_\epsilon^2}{2} \geq \mathrm{w}_\epsilon(\mathcal{M}_{s_\epsilon})$$

$$\text{where} \quad \mathrm{w}_\epsilon(\mathcal{V}) = \mathrm{E} \max_{v \in \mathcal{V}} \langle \epsilon, v \rangle_{L_2(\mathrm{P_n})} \quad \text{and} \quad \Sigma_n = \mathrm{E} \max_{i \in 1 \ldots n} \varepsilon_i^2.$$

Where we're going. We'll derive a bound that depends on the model $\mathcal{M}$ and the distribution of the noise ε in simpler and disentangled way.

$$\|\hat{\mu} - \mu^\star\|_{L_2(\mathrm{P_n})} < \sqrt{2\pi} M_n \, s + 2\sqrt{\frac{2\Sigma_n}{\delta n}} \leq \quad \text{w.p. } 1-\delta \text{ for } \quad \frac{s^2}{2} \geq \mathrm{w}(\mathcal{M}_s)$$

$$\text{where} \quad \mathrm{w}(\mathcal{V}) = \mathrm{E} \max_{v \in \mathcal{V}} \langle g, v \rangle_{L_2(\mathrm{P_n})} \quad \text{and} \quad M_n = \mathrm{E} \max_{i \in 1 \ldots n} |\varepsilon_i|.$$

Better yet. We can simplify it into a bound that depends on only one measure of noise complexity.

$$\|\hat{\mu} - \mu^\star\|_{L_2(\mathrm{P_n})} \leq \sqrt{2\pi\Sigma_n} \left(s + \sqrt{\frac{2}{\delta n}} \right) \quad \text{because} \quad M_n \leq \sqrt{\Sigma_n} \quad \text{and} \quad 2 \leq \sqrt{2\pi}$$

Non-Gaussian Noise

Example: Probabilistic Classification

The Setting

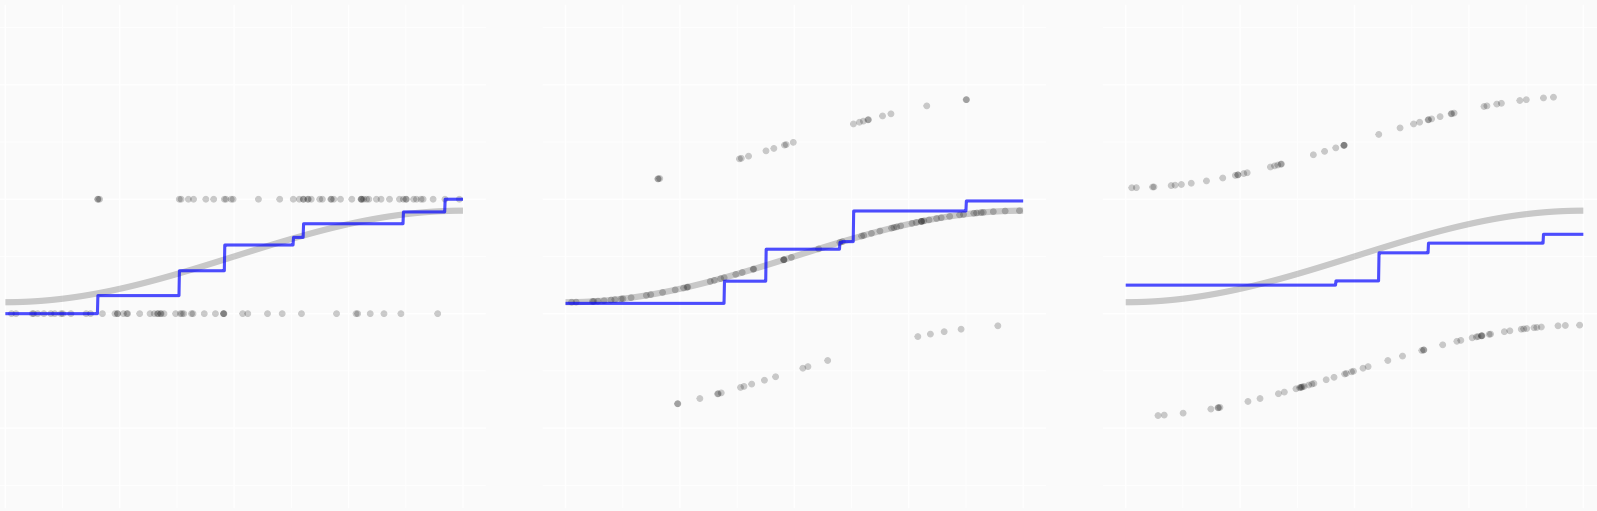

Figure 11: classification noise $\rightarrow$ symmetrized classification noise $\rightarrow$ random-sign noise

Suppose we have independent *binary observations*.

$$Y_i = \begin{cases} 1 & \text{with conditional probability } \mu(X_i) \\ 0 & \text{otherwise} \end{cases}$$

$$= \mu(X_i) + \varepsilon_i \quad \text{for} \quad \varepsilon_i = \begin{cases} 1 - \mu(X_i) & \text{with conditional probability } \mu(X_i) \\ -\mu(X_i) & \text{with conditional probability } 1 - \mu(X_i) \end{cases}.$$

Note that this *classification noise* ε_i has conditional mean zero.

$$\mathrm{E}[\varepsilon_i \mid X_i] = \mu(X_i)\{1 - \mu(X_i)\} + \{1 - \mu(X_i)\}\{-\mu(X_i)\} = 0.$$

The Setting

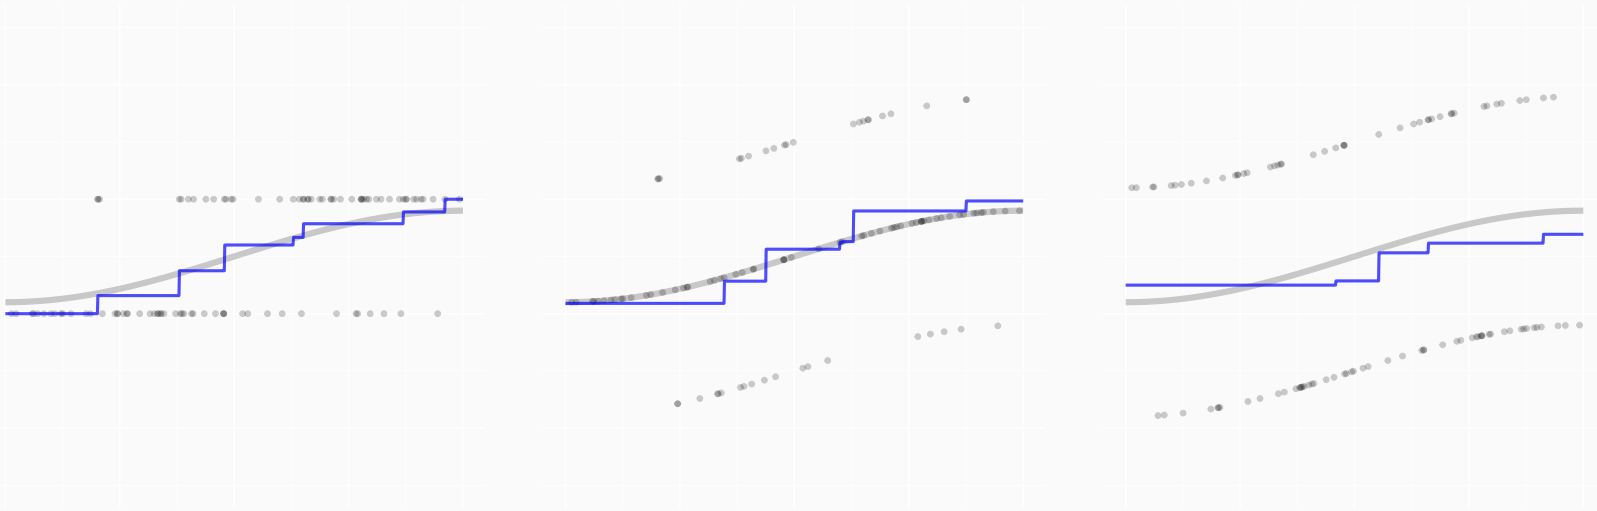

Figure 11: classification noise $\rightarrow$ symmetrized classification noise $\rightarrow$ random-sign noise

What we need to bound is *classification-noise width*

$$\mathrm{w}_\epsilon(\mathcal{V}) = \frac{1}{n}\,\mathrm{E}\max_{v\in\mathcal{V}}\sum_{i=1}^{n}\varepsilon_i v_i.$$

We'll show it's no bigger than a version with *symmetrized noise*.

$$\varepsilon_i - \varepsilon_i' = \begin{cases} +1 & \text{when } \varepsilon_i = 1-\mu(X_i),\ \varepsilon_i' = \mu(X_i) \\ -1 & \text{when } \varepsilon_i = \mu(X_i),\ \varepsilon_i' = 1-\mu(X_i) \\ 0 & \text{when } \varepsilon_i = \varepsilon_i' \end{cases}$$

The Setting

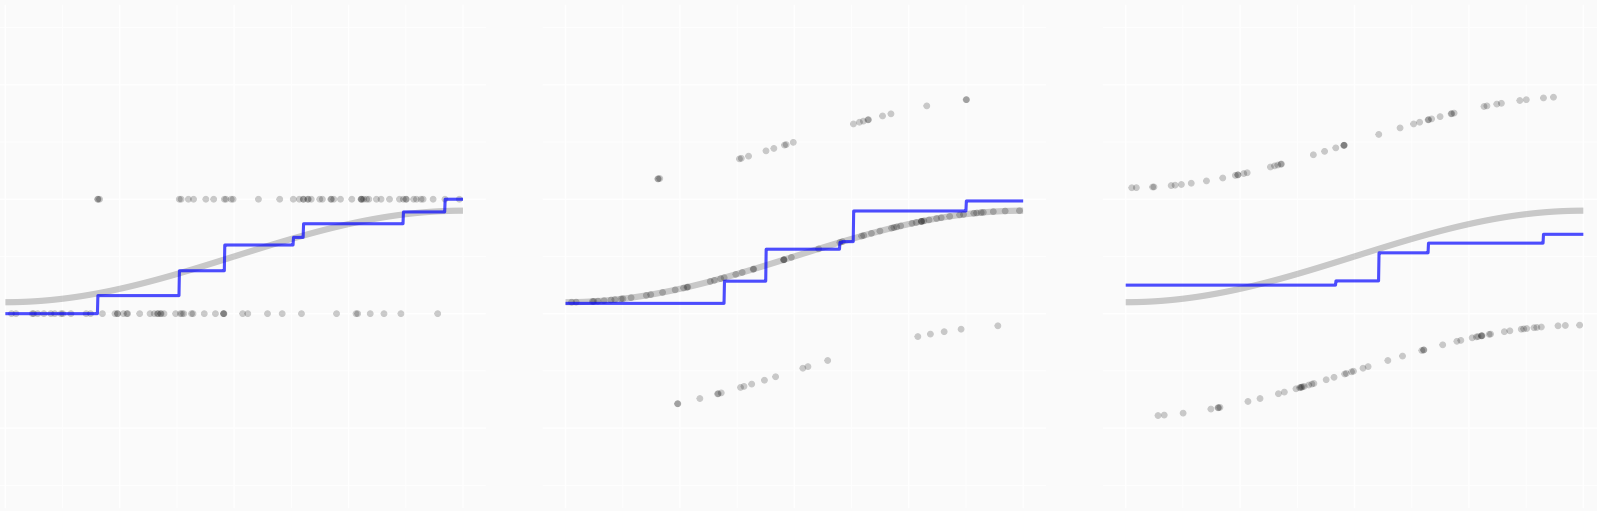

Figure 11: classification noise $\rightarrow$ symmetrized classification noise $\rightarrow$ random-sign noise

And we'll show that *this* is no bigger than a version with *random sign noise*

$$\mathrm{w}_\epsilon(\mathcal{V}) \leq \mathrm{w}_{\epsilon-\epsilon'}(\mathcal{V}) \leq \mathrm{w}_s(\mathcal{V}) \quad \text{where} \quad s_i = \pm 1 \text{ w.p. } 1/2.$$

The trick will be multiplying the symmetrized noise by a random sign.
It's already symmetric, so that doesn't change its distribution.

$$\varepsilon_i - \varepsilon_i' \overset{dist}{=} s_i(\varepsilon_i - \varepsilon_i')$$

Then we'll *contract out* the symmetrized noise, leaving the random sign. You'll see.

Step 1

We bound our maximum in terms of one involving symmetric noise.
We'll work with an *independent copy* ε' of our noise vector ε.

$$\begin{aligned}
\mathrm{E}_\varepsilon \max_{v\in\mathcal{V}} \sum_{i=1}^n \varepsilon_i v_i &\overset{(a)}{=} \mathrm{E}_\varepsilon \max_{v\in\mathcal{V}} \sum_{i=1}^n (\varepsilon_i - \mathrm{E}_{\varepsilon'}\, \varepsilon_i') v_i \\
&\overset{(b)}{=} \mathrm{E}_\varepsilon \max_{v\in\mathcal{V}}\ \mathrm{E}_{\varepsilon'} \sum_{i=1}^n (\varepsilon_i - \varepsilon_i') v_i \\
&\overset{(c)}{\leq} \mathrm{E}_\varepsilon\, \mathrm{E}_{\varepsilon'} \max_{v\in\mathcal{V}} \sum_{i=1}^n (\varepsilon_i - \varepsilon_i') v_i.
\end{aligned}$$

Why do these steps work?

(a) $\mathrm{E}_{\varepsilon'}\, \varepsilon_i' = 0.$

Step 1

We bound our maximum in terms of one involving symmetric noise.
We'll work with an *independent copy* ε' of our noise vector ε.

$$\begin{aligned}
\mathrm{E}_\varepsilon \max_{v \in \mathcal{V}} \sum_{i=1}^n \varepsilon_i v_i &\overset{(a)}{=} \mathrm{E}_\varepsilon \max_{v \in \mathcal{V}} \sum_{i=1}^n (\varepsilon_i - \mathrm{E}_{\varepsilon'} \varepsilon_i') v_i \\
&\overset{(b)}{=} \mathrm{E}_\varepsilon \max_{v \in \mathcal{V}} \mathrm{E}_{\varepsilon'} \sum_{i=1}^n (\varepsilon_i - \varepsilon_i') v_i \\
&\overset{(c)}{\leq} \mathrm{E}_\varepsilon \mathrm{E}_{\varepsilon'} \max_{v \in \mathcal{V}} \sum_{i=1}^n (\varepsilon_i - \varepsilon_i') v_i.
\end{aligned}$$

Why do these steps work?

(a) $\mathrm{E}_{\varepsilon'} \varepsilon_i' = 0$.

(b) Expectation is linear.

Step 1

We bound our maximum in terms of one involving symmetric noise.
We'll work with an *independent copy* ε' of our noise vector ε.

$$\begin{aligned}
\mathrm{E}_\varepsilon \max_{v \in \mathcal{V}} \sum_{i=1}^n \varepsilon_i v_i &\overset{(a)}{=} \mathrm{E}_\varepsilon \max_{v \in \mathcal{V}} \sum_{i=1}^n (\varepsilon_i - \mathrm{E}_{\varepsilon'} \varepsilon_i') v_i \\
&\overset{(b)}{=} \mathrm{E}_\varepsilon \max_{v \in \mathcal{V}} \mathrm{E}_{\varepsilon'} \sum_{i=1}^n (\varepsilon_i - \varepsilon_i') v_i \\
&\overset{(c)}{\leq} \mathrm{E}_\varepsilon \mathrm{E}_{\varepsilon'} \max_{v \in \mathcal{V}} \sum_{i=1}^n (\varepsilon_i - \varepsilon_i') v_i.
\end{aligned}$$

Why do these steps work?

(a) $\mathrm{E}_{\varepsilon'} \varepsilon_i' = 0$.
(b) Expectation is linear.
(c) Maximizing the average gives us something smaller than averaging the maxima.

Step 1

We bound our maximum in terms of one involving symmetric noise.
We'll work with an *independent copy* ε' of our noise vector ε.

$$\begin{aligned}
\mathrm{E}_\varepsilon \max_{v \in \mathcal{V}} \sum_{i=1}^n \varepsilon_i v_i &\overset{(a)}{=} \mathrm{E}_\varepsilon \max_{v \in \mathcal{V}} \sum_{i=1}^n (\varepsilon_i - \mathrm{E}_{\varepsilon'} \varepsilon_i') v_i \\
&\overset{(b)}{=} \mathrm{E}_\varepsilon \max_{v \in \mathcal{V}} \mathrm{E}_{\varepsilon'} \sum_{i=1}^n (\varepsilon_i - \varepsilon_i') v_i \\
&\overset{(c)}{\leq} \mathrm{E}_\varepsilon \mathrm{E}_{\varepsilon'} \max_{v \in \mathcal{V}} \sum_{i=1}^n (\varepsilon_i - \varepsilon_i') v_i.
\end{aligned}$$

Why do these steps work?

(a) $\mathrm{E}_{\varepsilon'} \varepsilon_i' = 0$.

(b) Expectation is linear.

(c) Maximizing the average gives us something smaller than averaging the maxima.

- In (c), we choose the maximizing $v \in \mathcal{V}$ *for each* ε'.

Step 1

We bound our maximum in terms of one involving symmetric noise.
We'll work with an *independent copy* ε' of our noise vector ε.

$$\begin{aligned}
\mathrm{E}_\varepsilon \max_{v \in \mathcal{V}} \sum_{i=1}^n \varepsilon_i v_i &\overset{(a)}{=} \mathrm{E}_\varepsilon \max_{v \in \mathcal{V}} \sum_{i=1}^n (\varepsilon_i - \mathrm{E}_{\varepsilon'} \varepsilon_i') v_i \\
&\overset{(b)}{=} \mathrm{E}_\varepsilon \max_{v \in \mathcal{V}} \mathrm{E}_{\varepsilon'} \sum_{i=1}^n (\varepsilon_i - \varepsilon_i') v_i \\
&\overset{(c)}{\leq} \mathrm{E}_\varepsilon \mathrm{E}_{\varepsilon'} \max_{v \in \mathcal{V}} \sum_{i=1}^n (\varepsilon_i - \varepsilon_i') v_i.
\end{aligned}$$

Why do these steps work?

(a) $\mathrm{E}_{\varepsilon'} \varepsilon_i' = 0$.

(b) Expectation is linear.

(c) Maximizing the average gives us something smaller than averaging the maxima.

- In (c), we choose the maximizing $v \in \mathcal{V}$ *for each* ε'.
- If we wanted to choose the same one each time, like we do in (b), we could.

Step 2

We introduce independent random signs $s_i = \pm 1$ w.p. $1/2$, changing nothing.

$$\mathrm{E}_{\varepsilon}\,\mathrm{E}_{\varepsilon'}\max_{v\in\mathcal{V}}\sum_{i=1}^{n}(\varepsilon_i-\varepsilon_i')v_i = \mathrm{E}_{s}\,\mathrm{E}_{\varepsilon}\,\mathrm{E}_{\varepsilon'}\max_{v\in\mathcal{V}}\sum_{i=1}^{n}s_i(\varepsilon_i-\varepsilon_i')v_i.$$

Why does this change nothing?

Step 2

We introduce independent random signs $s_i = \pm 1$ w.p. $1/2$, changing nothing.

$$\mathrm{E}_{\varepsilon}\,\mathrm{E}_{\varepsilon'}\max_{v\in\mathcal{V}}\sum_{i=1}^{n}(\varepsilon_i-\varepsilon_i')v_i = \mathrm{E}_{s}\,\mathrm{E}_{\varepsilon}\,\mathrm{E}_{\varepsilon'}\max_{v\in\mathcal{V}}\sum_{i=1}^{n}s_i(\varepsilon_i-\varepsilon_i')v_i.$$

Why does this change nothing?

- Because the inner mean $(\mathrm{E}_{\varepsilon}\,\mathrm{E}_{\varepsilon'})$ doesn't depend on the signs s_i.
- That's because ε_i and ε_i' have the same distribution.
- And this implies $(\varepsilon_i - \varepsilon_i')$ and $(\varepsilon_i' - \varepsilon) = -(\varepsilon_i - \varepsilon_i')$ do, too.

Step 3

We swap the order of our averages and think about the inner average as a *function* of our vector of symmetric noise.

$$\mathrm{E}_s\,\mathrm{E}_\varepsilon\,\mathrm{E}_{\varepsilon'}\max_{v\in\mathcal{V}}\sum_{i=1}^n s_i(\varepsilon_i-\varepsilon_i')v_i = \mathrm{E}_\varepsilon\,\mathrm{E}_{\varepsilon'}\,\mathrm{E}_s\max_{v\in\mathcal{V}}\sum_{i=1}^n s_i(\varepsilon_i-\varepsilon_i')v_i$$

$$= \mathrm{E}_\varepsilon\,\mathrm{E}_{\varepsilon'}\,f(\varepsilon-\varepsilon') \quad \text{for} \quad f(u) = \mathrm{E}_s\max_{v\in\mathcal{V}}\sum_{i=1}^n s_i u_i v_i.$$

This function f is convex.

What does that mean? These, for example, are all convex.

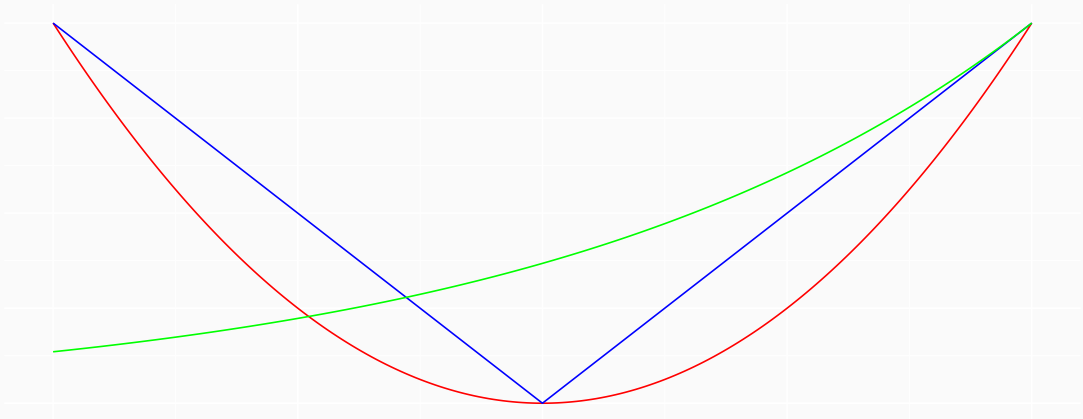

$f\{(1-\lambda)a+\lambda b\} \le (1-\lambda)f(a)+\lambda f(b) \quad$ for $\quad \lambda\in[0,1]$. That's Convexity

Step 3

We swap the order of our averages and think about the inner average as a *function* of our vector of symmetric noise.

$$\mathrm{E}_s \mathrm{E}_\varepsilon \mathrm{E}_{\varepsilon'} \max_{v \in \mathcal{V}} \sum_{i=1}^n s_i(\varepsilon_i - \varepsilon_i')v_i = \mathrm{E}_\varepsilon \mathrm{E}_{\varepsilon'} \mathrm{E}_s \max_{v \in \mathcal{V}} \sum_{i=1}^n s_i(\varepsilon_i - \varepsilon_i')v_i$$

$$= \mathrm{E}_\varepsilon \mathrm{E}_{\varepsilon'} f(\varepsilon - \varepsilon') \quad \text{for} \quad f(u) = \mathrm{E}_s \max_{v \in \mathcal{V}} \sum_{i=1}^n s_i u_i v_i.$$

This function f is convex.

How do we know? Maximizing each term is better than maximizing their sum.

$$f\{(1-\lambda)a + \lambda b\} = \mathrm{E}_s \max_{v \in \mathcal{V}} \left\{ (1-\lambda) \sum_{i=1}^n s_i a_i v_i + \lambda \sum_{i=1}^n s_i b_i v_i \right\}$$

$$\leq \mathrm{E}_s \left\{ \max_{v \in \mathcal{V}} (1-\lambda) \sum_{i=1}^n s_i a_i v_i + \max_{v \in \mathcal{V}} \lambda \sum_{i=1}^n s_i b_i v_i \right\}$$

$$= (1-\lambda) \mathrm{E}_s \max_{v \in \mathcal{V}} \sum_{i=1}^n s_i a_i v_i + \lambda \mathrm{E}_s \max_{v \in \mathcal{V}} \lambda \sum_{i=1}^n s_i b_i v_i$$

$$= (1-\lambda) f(a) + \lambda f(b).$$

Step 3

We swap the order of our averages and think about the inner average as a *function* of our vector of symmetric noise.

$$\mathrm{E}_s \mathrm{E}_\varepsilon \mathrm{E}_{\varepsilon'} \max_{v \in \mathcal{V}} \sum_{i=1}^{n} s_i(\varepsilon_i - \varepsilon_i')v_i = \mathrm{E}_\varepsilon \mathrm{E}_{\varepsilon'} \mathrm{E}_s \max_{v \in \mathcal{V}} \sum_{i=1}^{n} s_i(\varepsilon_i - \varepsilon_i')v_i$$

$$= \mathrm{E}_\varepsilon \mathrm{E}_{\varepsilon'} f(\varepsilon - \varepsilon') \quad \text{for} \quad f(u) = \mathrm{E}_s \max_{v \in \mathcal{V}} \sum_{i=1}^{n} s_i u_i v_i.$$

This function f is convex.

Why does this matter? The max of a convex function over a cube occurs at a corner.

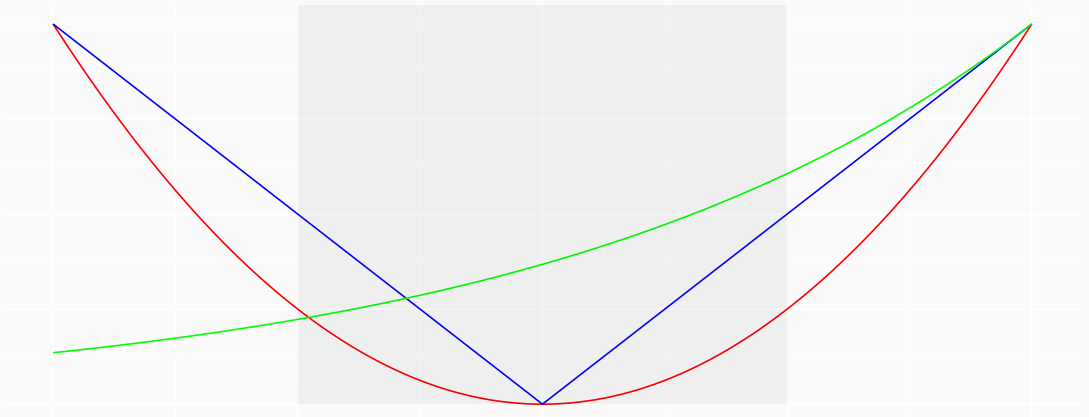

What cube?

Step 4

The vector of symmetric noise, $\varepsilon - \varepsilon'$, is in the *unit cube* $[-1, 1]^n$.

$$\varepsilon_i - \varepsilon'_i = \begin{cases} 0 & \text{when } \varepsilon_i = \varepsilon'_i \\ +1 & \text{when } \varepsilon_i = 1 - \mu(X_i),\ \varepsilon'_i = \mu(X_i) \\ -1 & \text{when } \varepsilon_i = \mu(X_i),\ \varepsilon'_i = 1 - \mu(X_i). \end{cases}$$

The average over this random vector is bounded by the maximum over the cube it's in.

$$\begin{aligned} \mathrm{E}_\varepsilon \, \mathrm{E}_{\varepsilon'} \, \mathrm{E}_s \max_{v \in \mathcal{V}} \sum_{i=1}^n s_i(\varepsilon_i - \varepsilon'_i) v_i &\leq \max_{u \in [-1,1]^n} \mathrm{E}_s \max_{v \in \mathcal{V}} \sum_{i=1}^n s_i u_i v_i \\ &= \max_{u \in [-1,1]^n} f(u) \qquad \text{max over the cube} \\ &= \max_{u \in \{-1,1\}^n} f(u) \qquad \text{max over its corners} \end{aligned}$$

Step 5

We characterize this maximum over corners. Remember what f is.

$$\max_{u \in \{-1,1\}^n} f(u) = \max_{u \in \{-1,1\}^n} \mathrm{E}_s \max_{v \in \mathcal{V}} \sum_{i=1}^{n} s_i u_i v_i$$

$$= \mathrm{E}_s \max_{v \in \mathcal{V}} \sum_{i=1}^{n} s_i v_i.$$

Why?

Hint. What's the distribution of s_i? And $s_i u_i$ for $u_i \in \{-1, 1\}$?

Step 5

We characterize this maximum over corners. Remember what f is.

$$\max_{u \in \{-1,1\}^n} f(u) = \max_{u \in \{-1,1\}^n} \mathrm{E}_s \max_{v \in \mathcal{V}} \sum_{i=1}^{n} s_i u_i v_i$$

$$= \mathrm{E}_s \max_{v \in \mathcal{V}} \sum_{i=1}^{n} s_i v_i.$$

Why?

Hint. What's the distribution of s_i? And $s_i u_i$ for $u_i \in \{-1, 1\}$?

- For $u_i \in \{-1, 1\}$, the distributions of u_i and $s_i u_i$ are the same.
- So the distribution of the sum, and its maximum, are the same at every corner u.
- Including the vector of all ones $u = (1, 1, \ldots, 1)$.

Summary

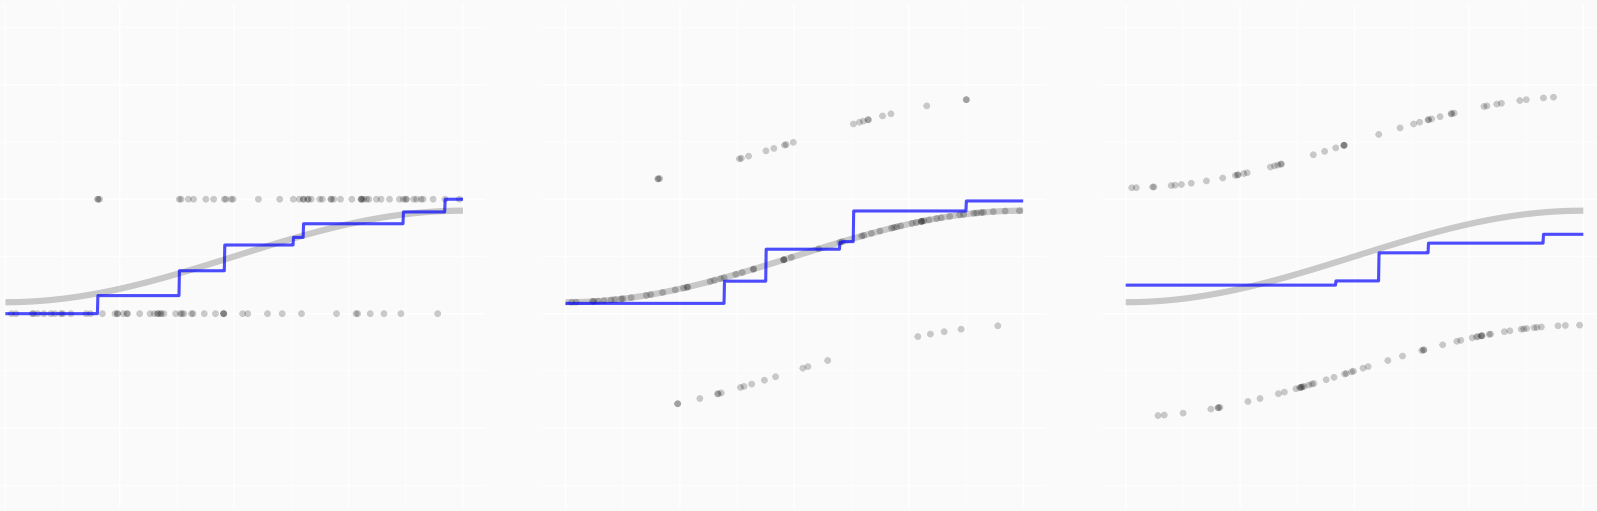

classification noise width $\leq$ symmetrized classification noise width $\leq$ random sign width

This means probabilistic classification is *easier* than regression
with random sign noise. Or, at least, that we get a better bound.

$$\frac{s^2}{2} \geq \mathrm{w}_s(\mathcal{M}_s) \quad \text{and} \quad \mathrm{w}_s(\mathcal{M}_s) \geq \mathrm{w}_\varepsilon(\mathcal{M}_s) \quad \Longrightarrow \quad \frac{s^2}{2} \geq \mathrm{w}_\varepsilon(\mathcal{M}_s)$$

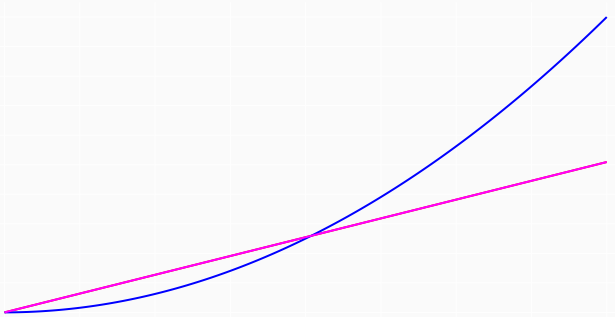

Terminology

People call random sign width, or something like it, *Rademacher Complexity*.

$$\underset{\mathrm{w}_s(\mathcal{V})}{\text{Rademacher Complexity}}(\mathcal{V}) = \mathbf{E} \max_{v \in \mathcal{V}} \langle s, v \rangle_{L_2(\mathbf{P_n})} \quad \text{for} \quad \text{i.i.d. } s_i = \begin{cases} +1 & \text{w.p. } 1/2 \\ -1 & \text{w.p. } 1/2 \end{cases}$$

$$\text{or maybe } = \mathbf{E} \max_{v \in \mathcal{V}} \left| \langle s, v \rangle_{L_2(\mathbf{P_n})} \right|$$

- This second definition is the same if $\mathcal{V}$ is symmetric, i.e. $v \in \mathcal{V} \implies -v \in \mathcal{V}$.
- Otherwise, it can be a little bigger.
  - At most $2\times$ bigger. Prove it!
  - Use the bound $\max a, b \leq a + b$ and the symmetry of s's distribution.

Non-Gaussian Noise

The General Case

Symmetrization and Contraction: Examples

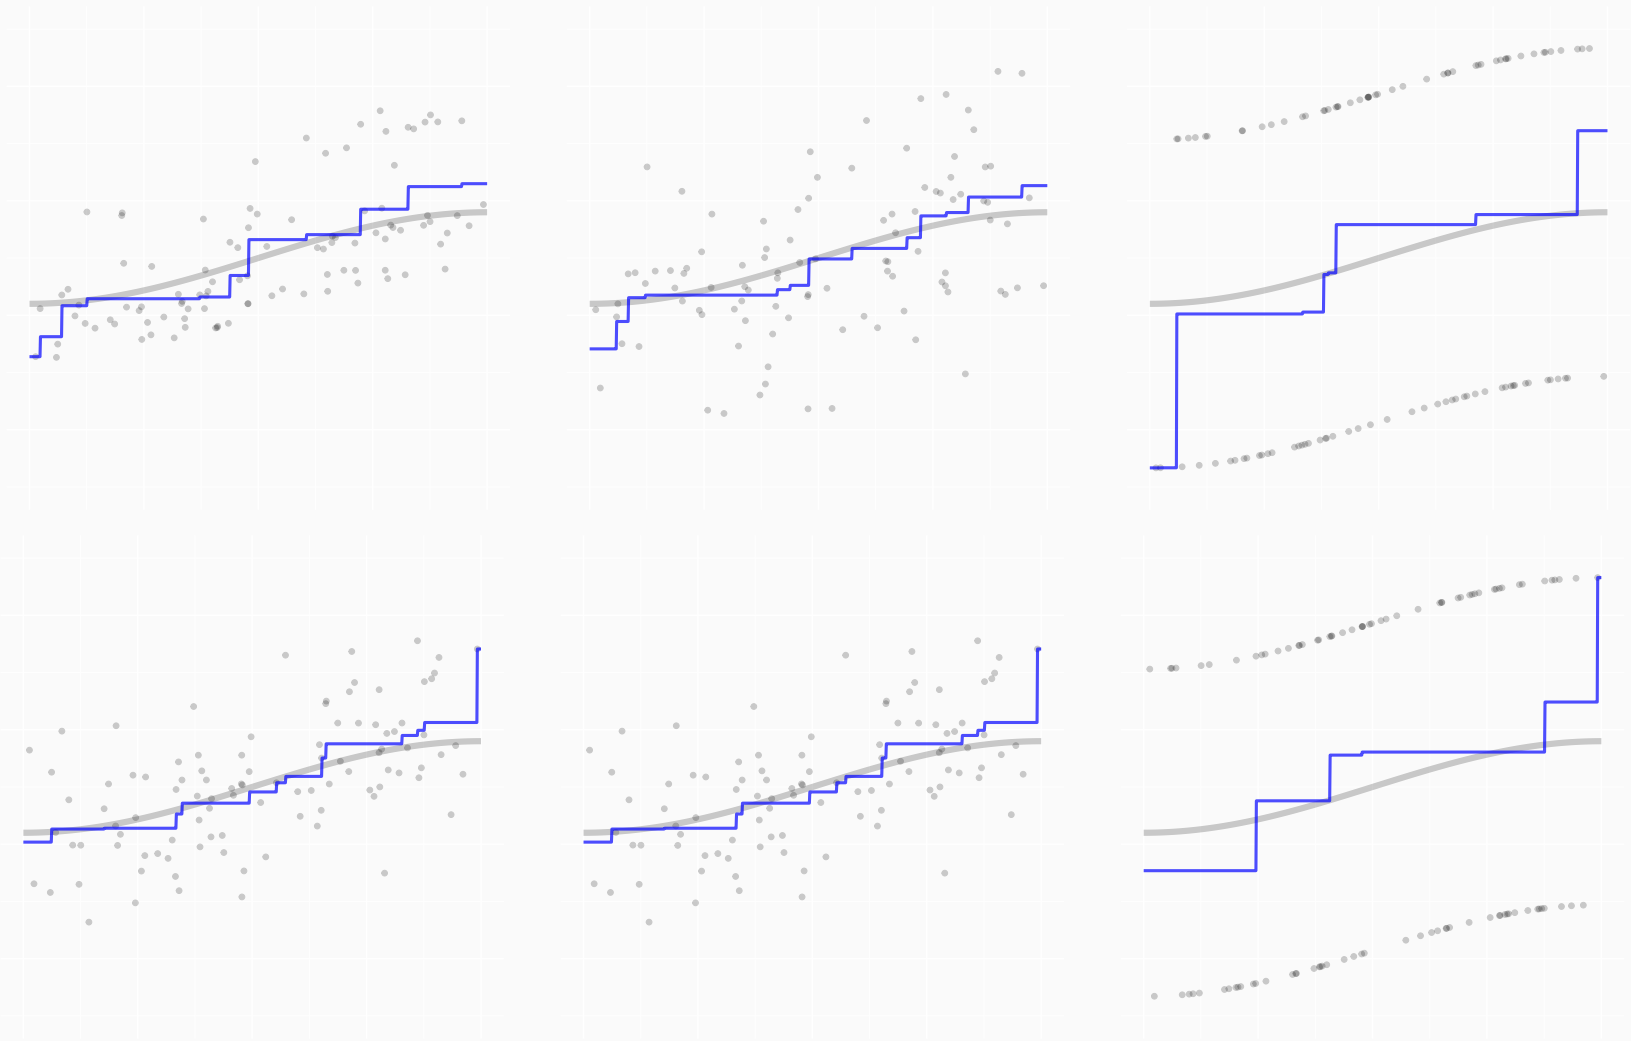

Figure 12: real noise $\longrightarrow$ symmetrized noise $\longrightarrow$ scaled sign noise

Symmetrization

$$\mathrm{w}_{\varepsilon}(\mathcal{V}) \le \mathrm{w}_{s(\varepsilon-\varepsilon')}(\mathcal{V}) \le 2\,\mathrm{w}_{s\varepsilon}(\mathcal{V})$$

$$\begin{aligned}
\mathrm{E} \max_{v\in\mathcal{V}} \sum_{i=1}^{n} \varepsilon_i v_i &= \mathrm{E} \max_{v\in\mathcal{V}} \sum_{i=1}^{n} (\varepsilon_i - \mathrm{E}\,\varepsilon_i') v_i \\
&\overset{(a)}{\le} \mathrm{E}\,\mathrm{E}' \max_{v\in\mathcal{V}} \sum_{i=1}^{n} (\varepsilon_i - \varepsilon_i') v_i \\
&= \mathrm{E}_s\,\mathrm{E}\,\mathrm{E}' \max_{v\in\mathcal{V}} \sum_{i=1}^{n} s_i(\varepsilon_i - \varepsilon_i') v_i \\
&\overset{(b)}{\le} \mathrm{E}_s\,\mathrm{E} \max_{v\in\mathcal{V}} \sum_{i=1}^{n} s_i\varepsilon_i + \mathrm{E}_s\,\mathrm{E}' \max_{v\in\mathcal{V}} \sum_{i=1}^{n} s_i\varepsilon_i' v_i \\
&= 2\,\mathrm{E}_s\,\mathrm{E} \max_{v\in\mathcal{V}} \sum_{i=1}^{n} \varepsilon_i s_i v_i.
\end{aligned}$$

(a) Replacing ε_i with $s_i(\varepsilon_i - \varepsilon_i')$ is 'free'.
- We stopped here in our example because $\varepsilon_i - \varepsilon_i'$ was easy to bound.
- Generally, we take an extra step to express things in terms of ε_i again.

(b) Replacing ε_i with $s_i\varepsilon_i$ increases width by at most $2\times$.

Contraction

$$\mathrm{w}_\eta(\mathcal{V}) = \mathrm{w}_{s\eta}(\mathcal{V}) \leq \mathrm{E}\|\eta\|_\infty \, \mathrm{w}_\eta(\mathcal{V}) \quad \text{if} \quad \eta \overset{dist}{=} -\eta.$$

$$\begin{aligned}
\mathrm{E}_s \, \mathrm{E}_\eta \max_{v \in \mathcal{V}} \sum_{i=1}^n \eta_i s_i v_i &\leq \mathrm{E}_\eta \max_{\substack{u \in \mathbb{R}^n \\ |u_i| \leq \|\eta\|_\infty}} \mathrm{E}_s \max_{v \in \mathcal{V}} \sum_{i=1}^n u_i s_i v_i \\
&= \mathrm{E}_\eta \|\eta\|_\infty \max_{u \in [-1,1]^n} \mathrm{E}_s \max_{v \in \mathcal{V}} \sum_{i=1}^n u_i s_i v_i \\
&= \mathrm{E}_\eta \|\eta\|_\infty \times \max_{u \in [-1,1]^n} \mathrm{E}_s \max_{v \in \mathcal{V}} \sum_{i=1}^n u_i s_i v_i \\
&= \mathrm{E}_\eta \|\eta\|_\infty \times \mathrm{E}_s \max_{v \in \mathcal{V}} \sum_{i=1}^n s_i v_i
\end{aligned}$$

- We can 'contract out' any symmetrically distributed noise vector η by ...
  1. multiplying in independent random signs s_i. Symmetry $\implies$ $s_i \eta_i \overset{dist}{=} \eta_i$.
  2. maximizing over a cube containing η.
- We just have to use a big enough cube.
  - In our example, $\eta = \varepsilon - \varepsilon'$ was in the unit cube $[-1, 1]^n$ deterministically.
  - Generally, we maximize over a random cube $[-\|\eta\|_\infty, \, \|\eta\|_\infty]^n$.
  - And we can pull out the cube's radius $\|\eta\|_\infty$ as a multiplicative factor.

Implications for Regression

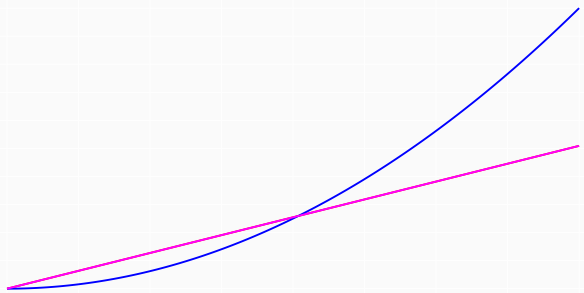

$$\mathrm{w}_\varepsilon(\mathcal{V}) \le \mathrm{E}\|\varepsilon_i - \varepsilon_i'\|_\infty \, \mathrm{w}_s(\mathcal{V}) \le 2\,\mathrm{E}\|\varepsilon_i\|_\infty \, \mathrm{w}_s(\mathcal{V})$$

Regression with arbitrary independent noise, i.e.

$$Y_i = \mu(X_i) + \varepsilon_i \quad \text{where} \quad \varepsilon_1 \dots \varepsilon_n \text{ are independent ,}$$

is no harder than with scaled-up random sign noise, i.e.

$$Y_i = \mu(X_i) + M s_i \quad \text{for} \quad M = \mathrm{E}\|\varepsilon_i - \varepsilon_i'\|_\infty \quad \text{and} \quad s_i = \begin{cases} +1 & \text{w.p. } 1/2 \\ -1 & \text{w.p. } 1/2 \end{cases}.$$

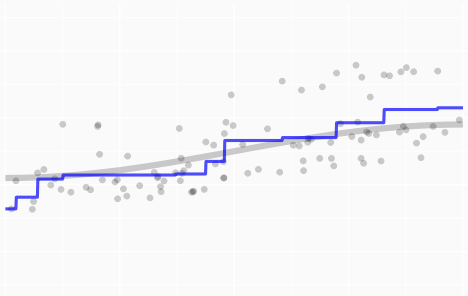

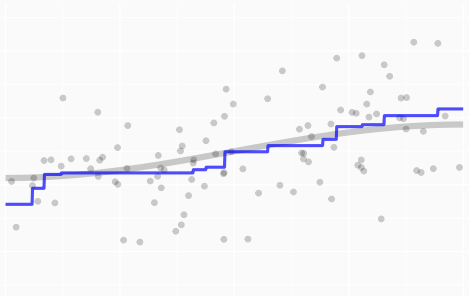

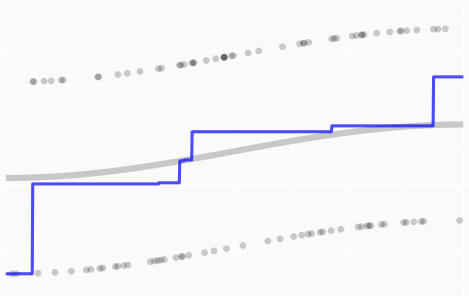

The Symmetric Case

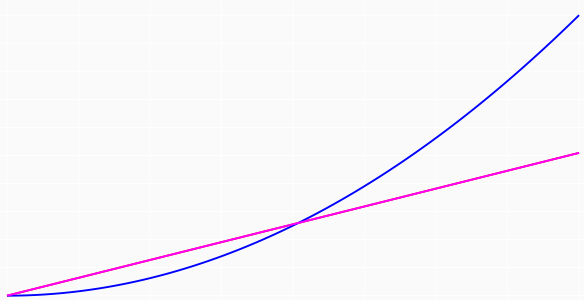

$$\mathrm{w}_\varepsilon(\mathcal{V}) \le \mathrm{E}\|\varepsilon_i\|_\infty \, \mathrm{w}_s(\mathcal{V})$$

Regression with arbitrary independent *symmetric* noise, i.e.

$$Y_i = \mu(X_i) + \varepsilon_i \quad \text{where} \quad \varepsilon_1 \dots \varepsilon_n \text{ are independent with } \varepsilon_i \overset{dist}{=} -\varepsilon_i,$$

is no harder than with scaled-up random sign noise, i.e.

$$Y_i = \mu(X_i) + M s_i \quad \text{for}^2 \quad M = \mathrm{E}\|\varepsilon_i\|_\infty \quad \text{and} \quad s_i = \begin{cases} +1 & \text{w.p. } 1/2 \\ -1 & \text{w.p. } 1/2 \end{cases}.$$

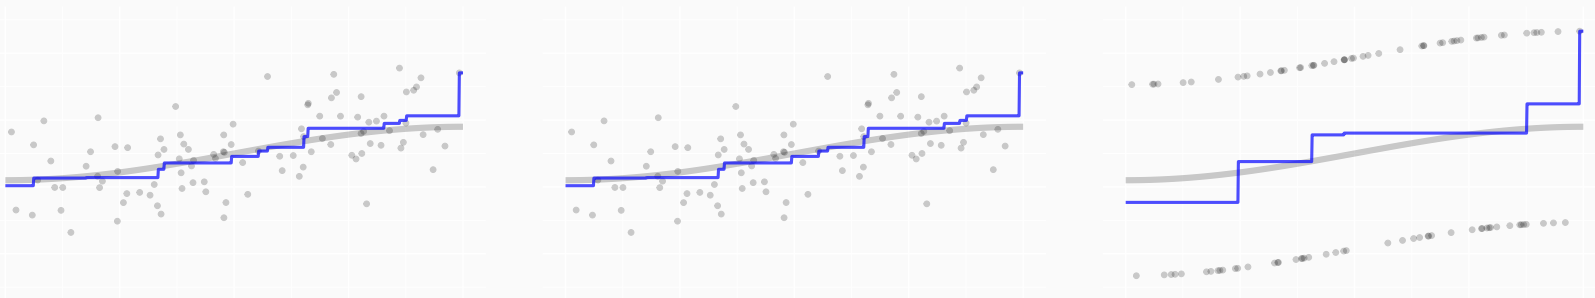

Figure 13: real noise $\to$ symmetrized noise $\to$ scaled sign noise

[2] $M = \mathrm{E}\|\varepsilon_i\|_\infty \le 2\sigma\sqrt{2\log(2n)}$ for $\varepsilon_i \sim N(0, \sigma^2)$. See Appendix B of the Gaussian Width Homework.

Non-Gaussian Noise

Comparison to the Gaussian Case

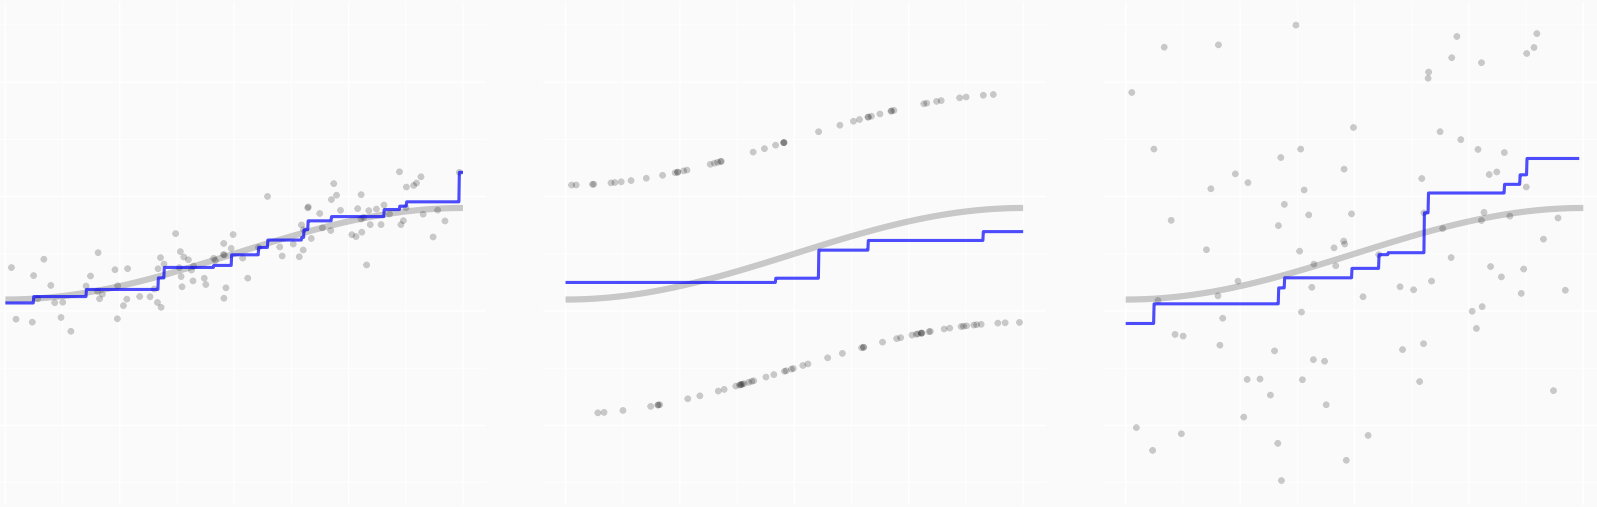

- So far, we've bounded arbitrary-noise width in terms of random-sign width.
- But often, it's easier to understand gaussian width. That's good enough.[3]

$$\underbrace{\frac{1}{2\sqrt{\log(2n)}}}_{\approx .2 \text{ for } n=100} \mathrm{w}_g(\mathcal{V}) \le \mathrm{w}_s(\mathcal{V}) \le \underbrace{\sqrt{\frac{\pi}{2}}}_{\approx 1.25} \mathrm{w}_g(\mathcal{V})$$

  - We just saw it can't be that much bigger than random-sign width.
  - And we can show it's at least $4/5$ as big.

$$\mathrm{E} \max_{v \in \mathcal{V}} \sum_{i=1}^{n} g_i v_i = \mathrm{E}_s \mathrm{E}_g \max_{v \in \mathcal{V}} \sum_{i=1}^{n} |g_i| s_i v_i \ge \mathrm{E}_s \max_{v \in \mathcal{V}} \sum_{i=1}^{n} \underbrace{\mathrm{E}_g |g_i|}_{=\sqrt{\frac{2}{\pi}}} s_i v_i.$$

[3] We can show $.125\, \mathrm{w}_g(\mathcal{V}) \le \mathrm{w}_s(\mathcal{V}) \le 1.25\, \mathrm{w}_g(\mathcal{V})$ for $n \le 10$ trillion by bounding $\mathrm{E}\|g\|_\infty$ more carefully.

Comparison in Steps

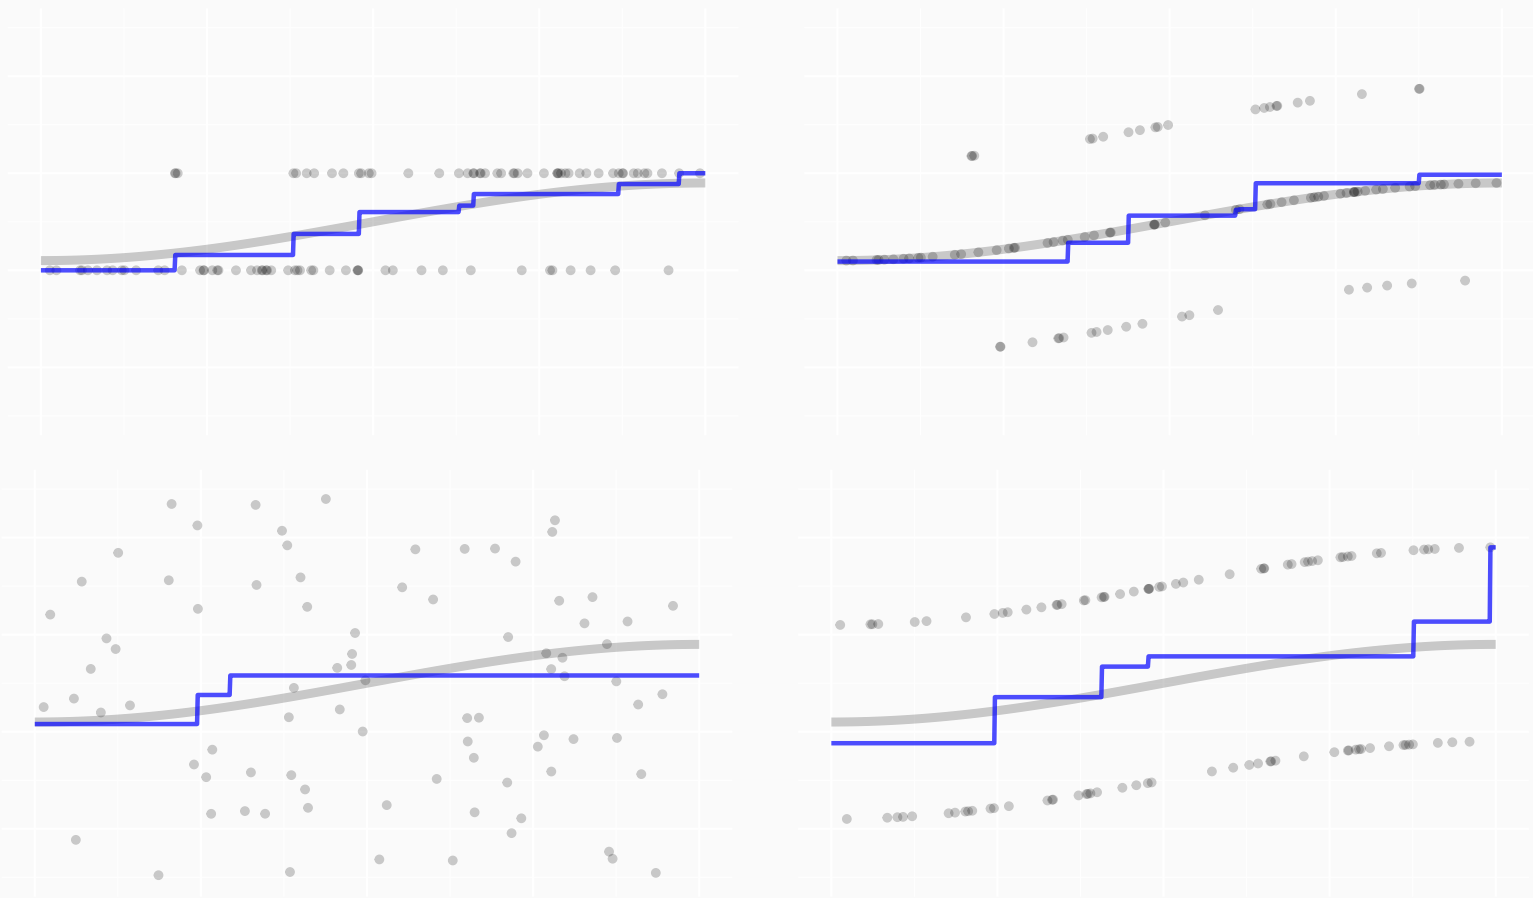

Figure 14: real noise → symmetrized noise ↓ scaled sign noise ← scaled gaussian noise

$$\mathrm{w}_\varepsilon(\mathcal{V}) \le \mathrm{w}_{\varepsilon-\varepsilon'}(\mathcal{V}) \le \underbrace{\mathrm{E}\|\varepsilon-\varepsilon'\|_\infty}_{\le 2\,\mathrm{E}\|\varepsilon\|_\infty} \quad \mathrm{w}_s(\mathcal{V}) \le \underbrace{\sqrt{\frac{\pi}{2}}\,\mathrm{E}\|\varepsilon-\varepsilon'\|_\infty}_{\le\sqrt{2\pi}\approx 2.5\times\mathrm{E}\|\varepsilon\|_\infty} \quad \mathrm{w}_g(\mathcal{V})$$

Implications for Regression

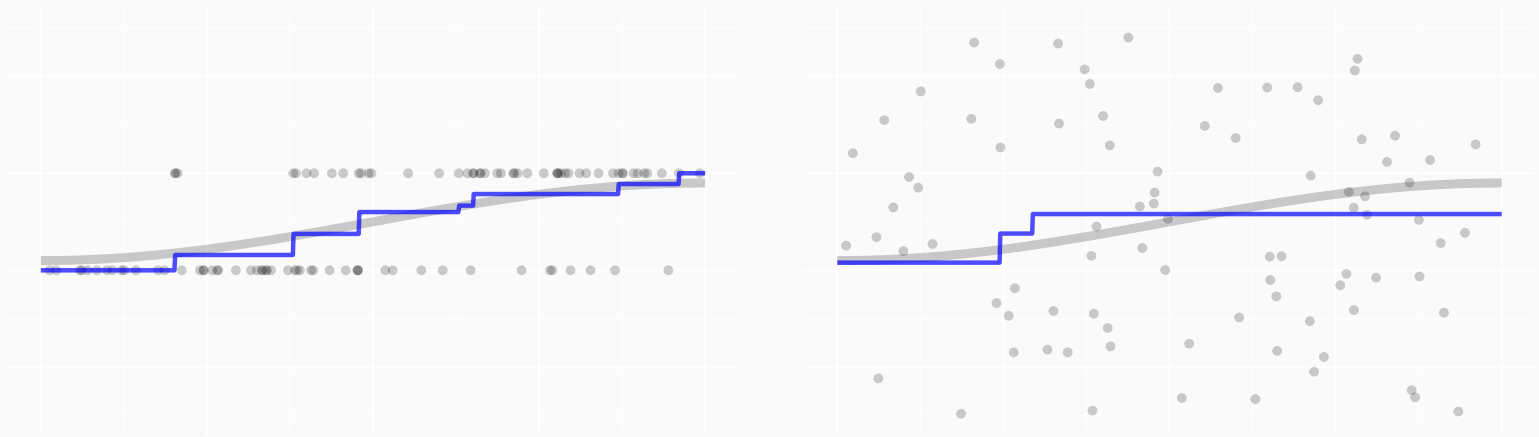

Figure 15: real noise $\to$ scaled gaussian noise

For any noise vector ε with independent components ε_i,

$$\mathrm{w}_\varepsilon(\mathcal{V}) \le 2\,\mathrm{E}\|\varepsilon\|_\infty \cdot \mathrm{w}_s(\mathcal{V}) \le \sqrt{2\pi}\,\mathrm{E}\|\varepsilon\|_\infty \cdot \mathrm{w}_g(\mathcal{V}).$$

- We can bound the width w_ε in terms of
  1. random-sign width
  2. the maximum absolute value of ε's components.
- And we can bound random-sign width in terms of gaussian width.

This means we don't have to bound a million different kinds of widths for each model.
We can bound random-sign width or gaussian width. Whichever is easier.

Sampling

Starting Point

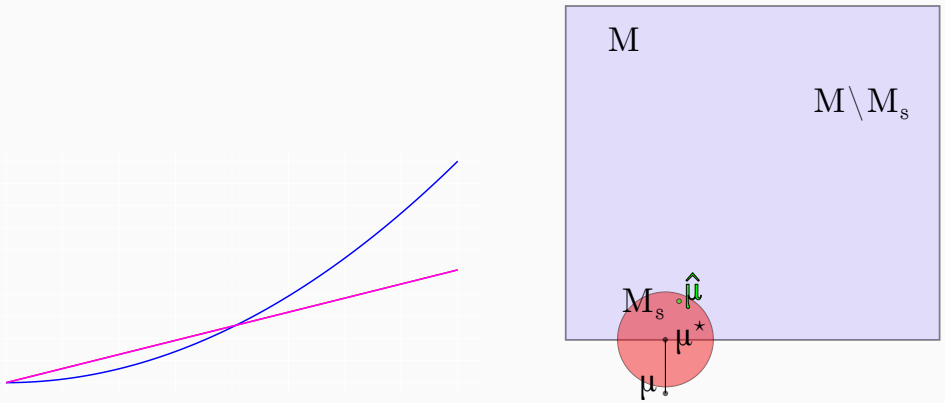

We have a bound that's valid for any signal μ and any vector of independent noise ε.

$$\|\hat{\mu} - \mu^\star\|_{L_2(\mathrm{P_n})} < \sqrt{2\pi\Sigma_n}\left(s + \sqrt{\frac{2}{\delta n}}\right) \quad \text{w.p. } 1-\delta$$

- It depends on the model's size through the *critical radius*.

  s satisfying $s^2/2 \geq \mathrm{w}(\mathcal{M}_s)$ for $\mathcal{M}_s = \{m \in \mathcal{M} \,:\, \|m - \mu^\star\|_{L_2(\mathrm{P_n})} \leq s\}$

  - This is a one-number summary of the gaussian width of neighborhoods ...
  - ...of the model's best approximation to the signal. It's the summary that matters.
- It depends on the noise's size through the expected maximum square.

$$\Sigma_n = \mathrm{E} \max_{i \in 1 \ldots n} |\varepsilon_i|^2$$

What does this tell us?

Bounds like this say how close $\hat{\mu}$ and $\mu^\star$ are, on average, on our sample $X_1 \dots X_n$.

$$\frac{1}{n}\sum_{i=1}^{n}\{\hat{\mu}(X_i) - \mu^\star(X_i)\} < \dots$$

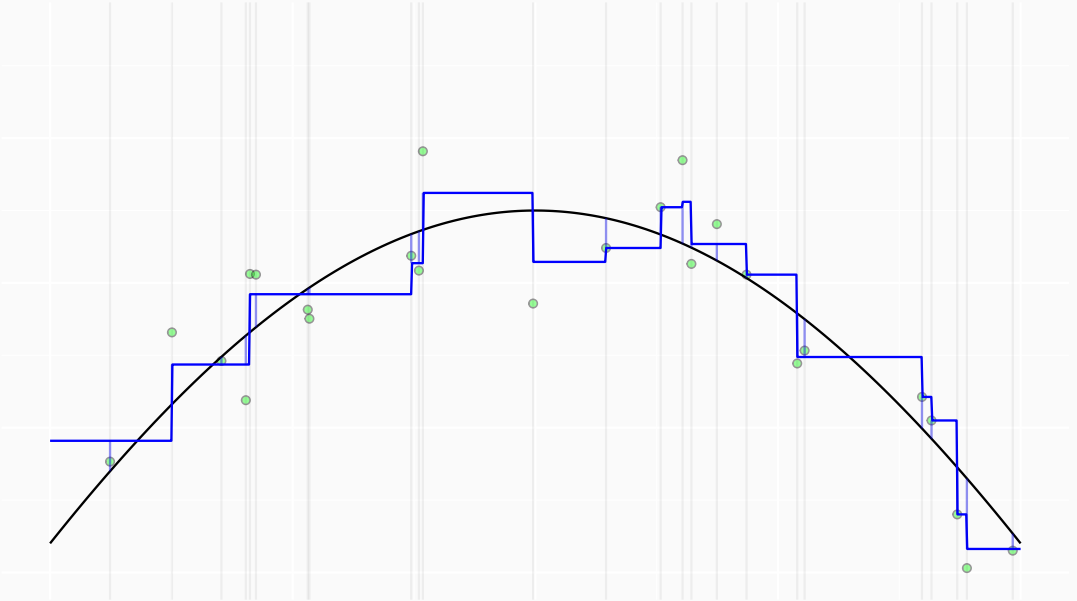

It doesn't tell us how close they are in the gaps between those points.

- Let's think about what happens when $X_1 \dots X_n$ is are drawn independently from some distribution P. Think sampling with replacement from a population.
- We'll bound the *population root mean squared error* $\|\hat{\mu} - \mu^\star\|_{L_2(\mathrm{P})}$.

What Population Mean Squared Error Is

It's the mean squared error we make at random point X' distributed like $X_1 \dots X_n$.

$$\|\hat{\mu} - \mu^\star\|^2_{L_2(\mathrm{P})} = \mathrm{E}_{X'}\left[\{\hat{\mu}(X') - \mu^\star(X')\}^2\right]$$

That's the integral of the squared distance between the two curves, multiplied by the density of X_i.

$$\|\hat{\mu} - \mu^\star\|^2_{L_2(\mathrm{P})} = \int \{\hat{\mu}(x) - \mu^\star(x)\}^2 p(x)\, dx \quad \text{if} \quad X_i \quad \text{has the density} \quad p(x).$$

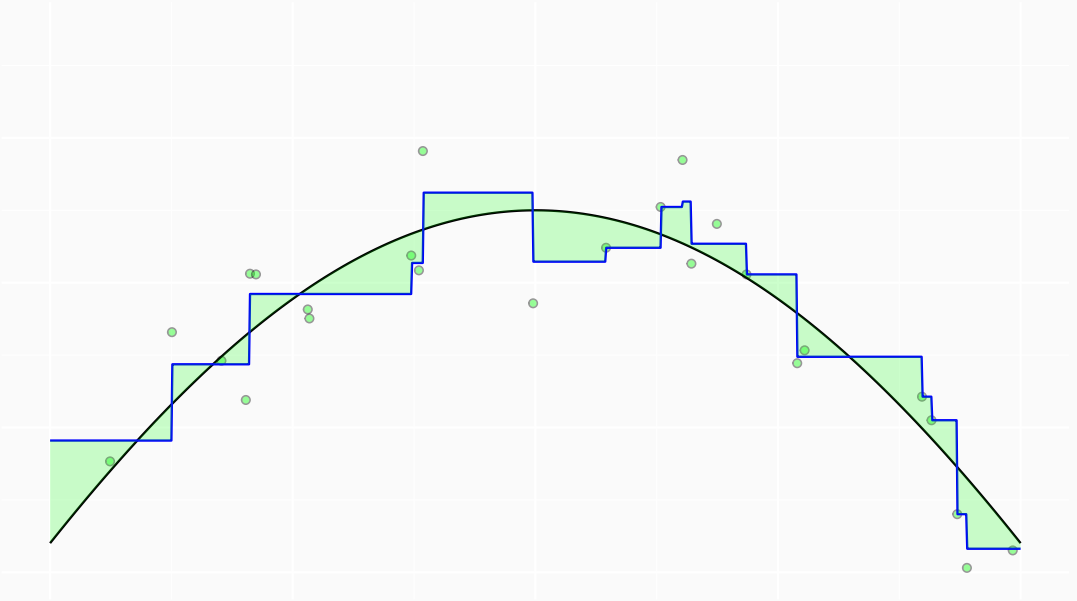

Why we care about Population Mean Squared Error: Generalization

If we're interested in average accuracy for a bunch of new points $X'_1 \dots X'_{n'}$ distributed like $X_1 \dots X_n$, that's more or less exactly what it is.

$$\|\hat{\mu} - \mu\|^2_{L_2(\mathrm{P})} = \mathrm{E}_{X'}\left[\{\hat{\mu}(X') - \mu(X')\}^2\right] \overset{LLN}{\approx} \frac{1}{n'}\sum_{i=1}^{n'}\{\hat{\mu}(X'_i) - \mu(X'_i)\}^2.$$

This can be a bit different from accuracy on our original sample $X_1 \dots X_n$.

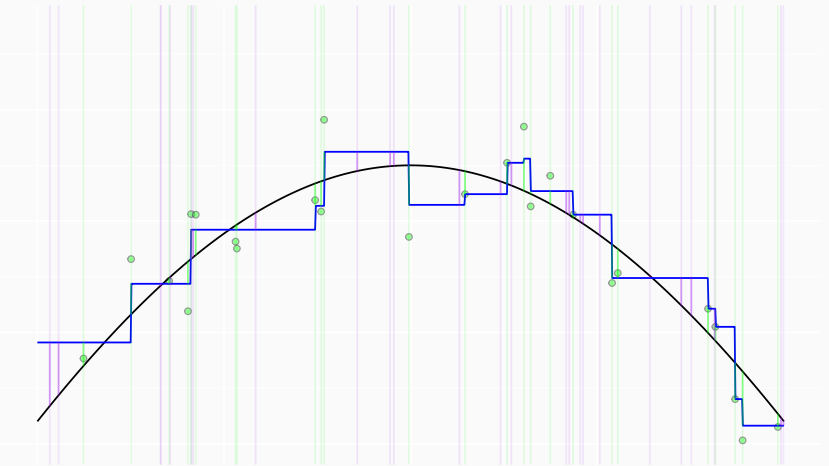

- BV regression spends its 'variation budget' jumping to fit on the original sample.
- Between those points, it doesn't know whether it should jump or not.
  - So we can get larger error at our new points.
  - It's usually not much larger, but sometimes it is. We'll see why.

Why we care about Population Mean Squared Error: Generalization

If we're interested in average accuracy for new points from a different distribution Q, we can bound this by comparing this distribution's density to that of our observations.

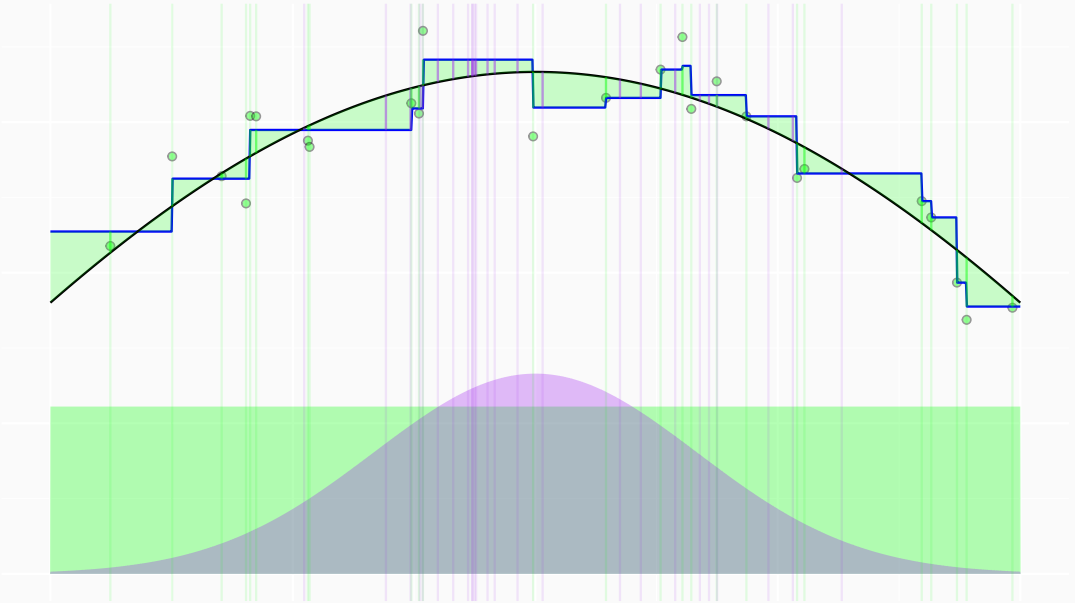

$$\frac{1}{n'}\sum_{i=1}^{n'}\{\hat{\mu}(X_i') - \mu(X_i')\}^2 \approx \|\hat{\mu} - \mu\|^2_{L_2(\mathrm{Q})} = \int \{\hat{\mu}(x) - \mu(x)\}^2 \frac{q(x)}{p(x)} p(x)\, dx$$

$$\leq \max_x \frac{q(x)}{p(x)} \|\hat{\mu} - \mu\|^2_{L_2(\mathrm{P})}.$$

Why we care about Population Mean Squared Error: Generalization

If we're interested in accuracy at a specific point x', we can think of this new distribution Q as a little bump around x'.

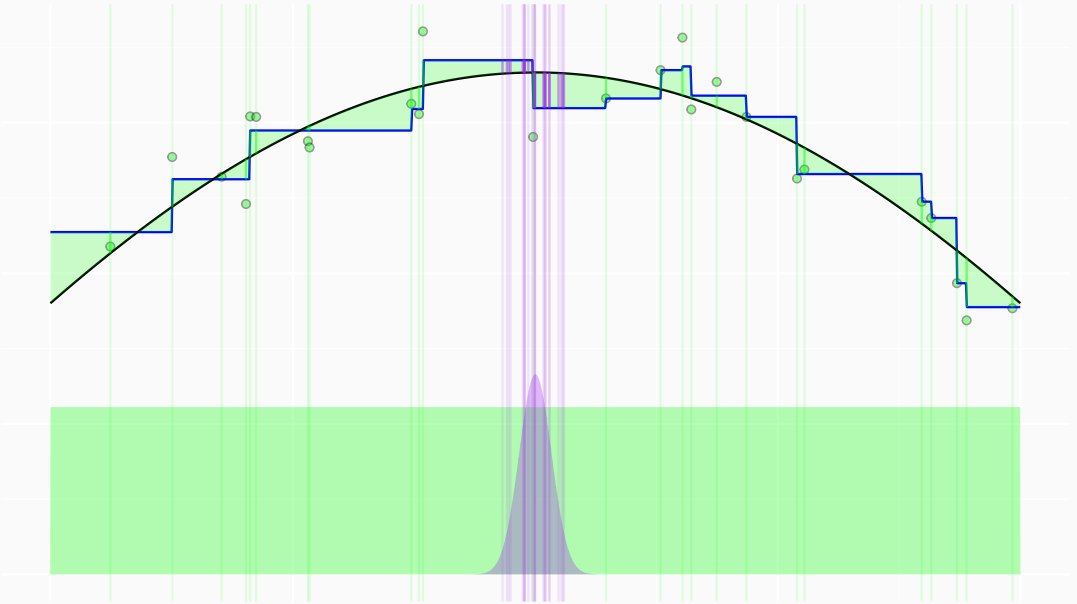

$$\{\hat{\mu}(x') - \mu(x')\}^2 \approx \|\hat{\mu} - \mu\|_{L_2(Q_\epsilon)} \quad \text{for} \quad Q = N(x', \epsilon^2).$$

What Does Our Bound Tell Us About This?

Idea. Suppose we know that $\hat{\mu}$ is in a neighborhood $\mathcal{M}_s$ of $\mu^\star$. We can ...

(a) Bound the expected maximum difference between population and sample mean squared error for curves in $\mathcal{M}_s$.
(b) Write population mean squared error as an expectation over an independent copy of our sample.
(c) Compare the result to a maximum of an average of symmetric random variables.

$$\begin{aligned}
\mathrm{E}\left[\|\hat{\mu}-\mu^\star\|_{L_2(\mathrm{P})}-\|\hat{\mu}-\mu^\star\|_{L_2(\mathrm{P_n})}\right] &\overset{(a)}{\leq} \mathrm{E}\max_{\delta\in\mathcal{M}_s-\mu^\star}\left|\|\delta\|_{L_2(\mathrm{P})}-\|\delta\|_{L_2(\mathrm{P_n})}\right| \\
&\overset{(b)}{=} \mathrm{E}_X \max_{\delta\in\mathcal{M}_s-\mu^\star} \mathrm{E}_{X'}\,\frac{1}{n}\sum_{i=1}^n \delta(X_i')^2-\frac{1}{n}\sum_{i=1}^n\delta(X_i)^2 \\
&\overset{(c)}{\leq} \mathrm{E}_X\,\mathrm{E}_{X'} \max_{\delta\in\mathcal{M}_s-\mu^\star}\frac{1}{n}\sum_{i=1}^n \delta(X_i')^2-\delta(X_i)^2
\end{aligned}$$

Now, in the spirit of our noise comparisons, we can try to compare this to something we know: the random sign width of $\mathcal{M}_s$. How does ...

$$\ldots = \mathrm{E}_X\,\mathrm{E}_{X'}\,\mathrm{E}_s \max_\delta \frac{1}{n}\sum^n s_i\{\delta(X_i')^2-\delta(X_i)^2\}$$

Contracting Out Lipschitz Functions

Claim. If ψ is a function with $\psi(0) = 0$ and $\rho_{\text{Lip}}(\psi) \leq L$, then

$$\mathrm{E}_s \max_f \sum_{i=1}^n s_i \psi_i\{f(X_i')\}^2 \leq L \max_f \sum_{i=1}^n s_i \delta(X_i)^2$$

Consequence. Letting $\psi(x) = x^2$,

$$\mathrm{E}_s \max_\delta \sum_{i=1}^n s_i \delta(X_i)^2 \leq L \max_\delta \sum_{i=1}^n s_i \delta(X_i) \quad \text{for} \quad L = \max_i |\psi'\{\delta(X_i)\}| = 2\|\delta\|_{L_\infty(\mathrm{P_n})}.$$

i.e.

$$\mathrm{E}_s \max_\delta \frac{1}{n} \sum_{i=1}^n s_i \delta(X_i)^2 \leq 2\|\delta\|_{L_\infty(\mathrm{P_n})}\, \mathrm{w}(\mathcal{M}_s)$$

Clever about it the infinity norm. 1. Bake in the constraint. 2. Bounded data — argue that you won't go above or below [e.g. monotone or lipschitz] 3. Suppose we know $\hat{\mu}$ *crosses* μ somewhere on $[0, 1]$. Then a derivative bound implies an infinity norm bound.